

Lecture 14: Interpreters

Eelco Visser

CS4200 Compiler Construction

TU Delft

December 2018

Reading Material

Design and implementation of DynSem, a domain-specific language for the definition of the dynamic semantics of programming languages.

DynSem provides facilities for implicitly passing memory (environment, store).

DynSem specifications can be compiled to or interpreted as interpreters for the languages they define.

RTA 2015

<http://dx.doi.org/10.4230/LIPICs.RTA.2015.365>

DynSem: A DSL for Dynamic Semantics Specification

Vlad Vergu, Pierre Neron, and Eelco Visser

Delft University of Technology
Delft, The Netherlands
{v.a.vergu|p.j.m.neron|visser}@tudelft.nl

Abstract

The formal semantics of a programming language and its implementation are typically separately defined, with the risk of divergence such that properties of the formal semantics are not properties of the implementation. In this paper, we present DynSem, a domain-specific language for the specification of the dynamic semantics of programming languages that aims at supporting both formal reasoning and efficient interpretation. DynSem supports the specification of the *operational semantics* of a language by means of *statically typed conditional term reduction rules*. DynSem supports *concise* specification of reduction rules by providing *implicit build and match coercions* based on reduction arrows and implicit term constructors. DynSem supports *modular* specification by adopting implicit propagation of semantic components from I-MSOS, which allows omitting propagation of components such as environments and stores from rules that do not affect those. DynSem supports the declaration of *native operators* for delegation of aspects of the semantics to an external definition or implementation. DynSem supports the definition of auxiliary *meta-functions*, which can be expressed using regular reduction rules and are subject to semantic component propagation. DynSem specifications are *executable* through automatic generation of a Java-based AST interpreter.

1998 ACM Subject Classification F.3.2. Semantics of Programming Languages (D.3.1.)

Keywords and phrases programming languages, dynamic semantics, reduction semantics, semantics engineering, IDE, interpreters, modularity

Digital Object Identifier 10.4230/LIPICs.RTA.2015.365

1 Introduction

The specification of the dynamic semantics is the core of a programming language design as it describes the runtime behavior of programs. In practice, the implementation of a compiler or an interpreter for the language often stands as the *only* definition of the semantics of a language. Such implementations, in a traditional programming language, often lack the clarity and the conciseness that a specification in a formal semantics framework provides. Therefore, they are a poor source of *documentation* about the semantics. On the other hand, formal definitions are not executable to the point that they can be used as implementations to run programs. Even when both a formal specification and an implementation co-exist, they typically diverge. As a result, important properties of a language as established based on its formal semantics may not hold for its implementation. Our goal is to unify the semantics engineering and language engineering of programming language designs [22] by providing a notation for the specification of the dynamic semantics that can serve both as a readable formalization as well as the source of an execution engine.

In this paper, we present DynSem, a DSL for the concise, modular, statically typed, and executable specification of the dynamic semantics of programming languages. DynSem

This paper describes the partial evaluation of the DynSem meta-interpreter with an object program.

The implementation makes use of Truffle a framework for run-time (online) partial evaluation.

Truffle gets most of it benefit from Graal, an implementation of the JVM with special support for partial evaluation.

ManLang 2018

<https://doi.org/10.1145/3237009.3237018>

Specializing a Meta-Interpreter

JIT Compilation of DynSem Specifications on the Graal VM

Vlad Vergu
TU Delft
The Netherlands
v.a.vergu@tudelft.nl

Eelco Visser
TU Delft
The Netherlands
visser@acm.org

ABSTRACT

DynSem is a domain-specific language for concise specification of the dynamic semantics of programming languages, aimed at rapid experimentation and evolution of language designs. DynSem specifications can be executed to interpret programs in the language under development. To enable fast turnaround during language development, we have developed a meta-interpreter for DynSem specifications, which requires minimal processing of the specification. In addition to fast development time, we also aim to achieve fast run times for interpreted programs.

In this paper we present the design of a meta-interpreter for DynSem and report on experiments with JIT compiling the application of the meta-interpreter on the Graal VM. By interpreting specifications directly, we have minimal compilation overhead. By specializing pattern matches, maintaining call-site dispatch chains and using native control-flow constructs we gain significant run-time performance. We evaluate the performance of the meta-interpreter when applied to the Tiger language specification running a set of common benchmark programs. Specialization enables the Graal VM to JIT compile the meta-interpreter giving speedups of up to factor 15 over running on the standard Oracle Java VM.

CCS CONCEPTS

• **Software and its engineering** → **Interpreters; Domain specific languages; Semantics;**

KEYWORDS

dynamic semantics, interpretation, JIT, run-time optimization

ACM Reference Format:

Vlad Vergu and Eelco Visser. 2018. Specializing a Meta-Interpreter: JIT Compilation of DynSem Specifications on the Graal VM. In *15th International Conference on Managed Languages & Runtimes (ManLang’18)*, September 12–14, 2018, Linz, Austria. ACM, New York, NY, USA, 14 pages. <https://doi.org/10.1145/3237009.3237018>

1 INTRODUCTION

The dynamic semantics of a programming language defines the run time execution behavior of programs in the language. Ideally,

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

ManLang’18, September 12–14, 2018, Linz, Austria

© 2018 Association for Computing Machinery.

ACM ISBN 978-1-4503-6424-9/18/09... \$15.00

<https://doi.org/10.1145/3237009.3237018>

the design of a programming language starts with the specification of its dynamic semantics to provide a high-level readable and unambiguous definition. However, understanding the design of a programming language also requires experimentation by actually running programs. Therefore, this ideal route is rarely taken, but language designs are embodied in the implementation of interpreters or compilers instead.

We have previously designed DynSem [33], a high-level meta-DSL for dynamic semantics specifications of programming languages, with the aim of supporting readable *and* executable specification. It supports the definition of modular and concise semantics by means of reduction rules with implicit propagation of contextual information. DynSem’s executable semantics entails that specifications can be used to interpret object language programs.

In our early prototypes, DynSem specifications were compiled to an interpreter. The process of generating a Java implementation of an interpreter and compiling that generated code caused long turnaround times during language prototyping. In order to support rapid prototyping with short turnaround times, we turned to interpreting specifications directly instead of compiling them. A DynSem interpreter is a *meta-interpreter* since the programs it interprets are themselves interpreters. Figure 1 depicts the high-level architecture of the DynSem meta-interpreter. First, a DynSem specification is desugared (explicated) to make implicit passing of semantic components explicit. The resulting specification in DynSem Core is then loaded into the meta-interpreter together with the AST of the interpreted object program. The interpreter consumes the program as input enacting the specification. This produces the desired result of a short turnaround time for experimenting with dynamic semantics specifications.


Meta-interpretation reduces the turnaround time at the expense of execution performance. At run time there are two interpreter layers operating (the meta-language interpreter and the object-language interpreter) which introduces substantial overhead. While we envision DynSem as a convenient way to prototype the dynamic semantics of programming languages, ultimately we also envision it as a convenient way to bridge the gap between the prototyping and production phases of a programming language’s lifecycle. Thus, we not only want an interpreter fast, but we also want a fast interpreter, which raises the question: Can we achieve fast object-language interpreters by optimizing the meta-interpretation of dynamic semantics specifications?

Direct vanilla interpreters are in general slow to begin with, even when they are implemented in a host language that is JIT-ed. This is because the host JIT is unable to see patterns in the object language and to meaningfully optimize the interpreter. The task of optimizing an interpreter has traditionally been long and

Semantics

What is the meaning of a program?

$$\text{meaning}(p) = \text{behavior}(p)$$



$\text{meaning}(p)$ = what happens when executing the generated (byte) code to which p is compiled

What is the meaning of a program?

$$\text{meaning}(p) = \text{behavior}(p)$$

What *is* behavior?

How can we *observe* behavior?

Mapping input to output

Changes to state of the system

Which behavior is essential, which accidental?

How can we define the semantics of a program?


Compiler defines *translational* semantics

$$\text{semanticsL1}(p) = \text{semanticsL2}(\text{translate}(p))$$

Requires understanding `translate` and `semanticsL2`



How do we know that `translate` is correct?

Is there a more ***direct description*** of `semanticsL1`?





**source
code**


semanticsL1





value



**machine
code**



semanticsL2



value

Verifying Compiler Correctness

Direct semantics of source language provides a specification

How to check correctness?

Testing: for *many* programs p (and inputs i) **test** that

$$\text{run}(p)(i) == \text{run}(\text{translate}(p))(i)$$

Verification: for *all* programs p (and inputs i) **prove** that

$$\text{run}(p)(i) == \text{run}(\text{translate}(p))(i)$$

Validating Semantics

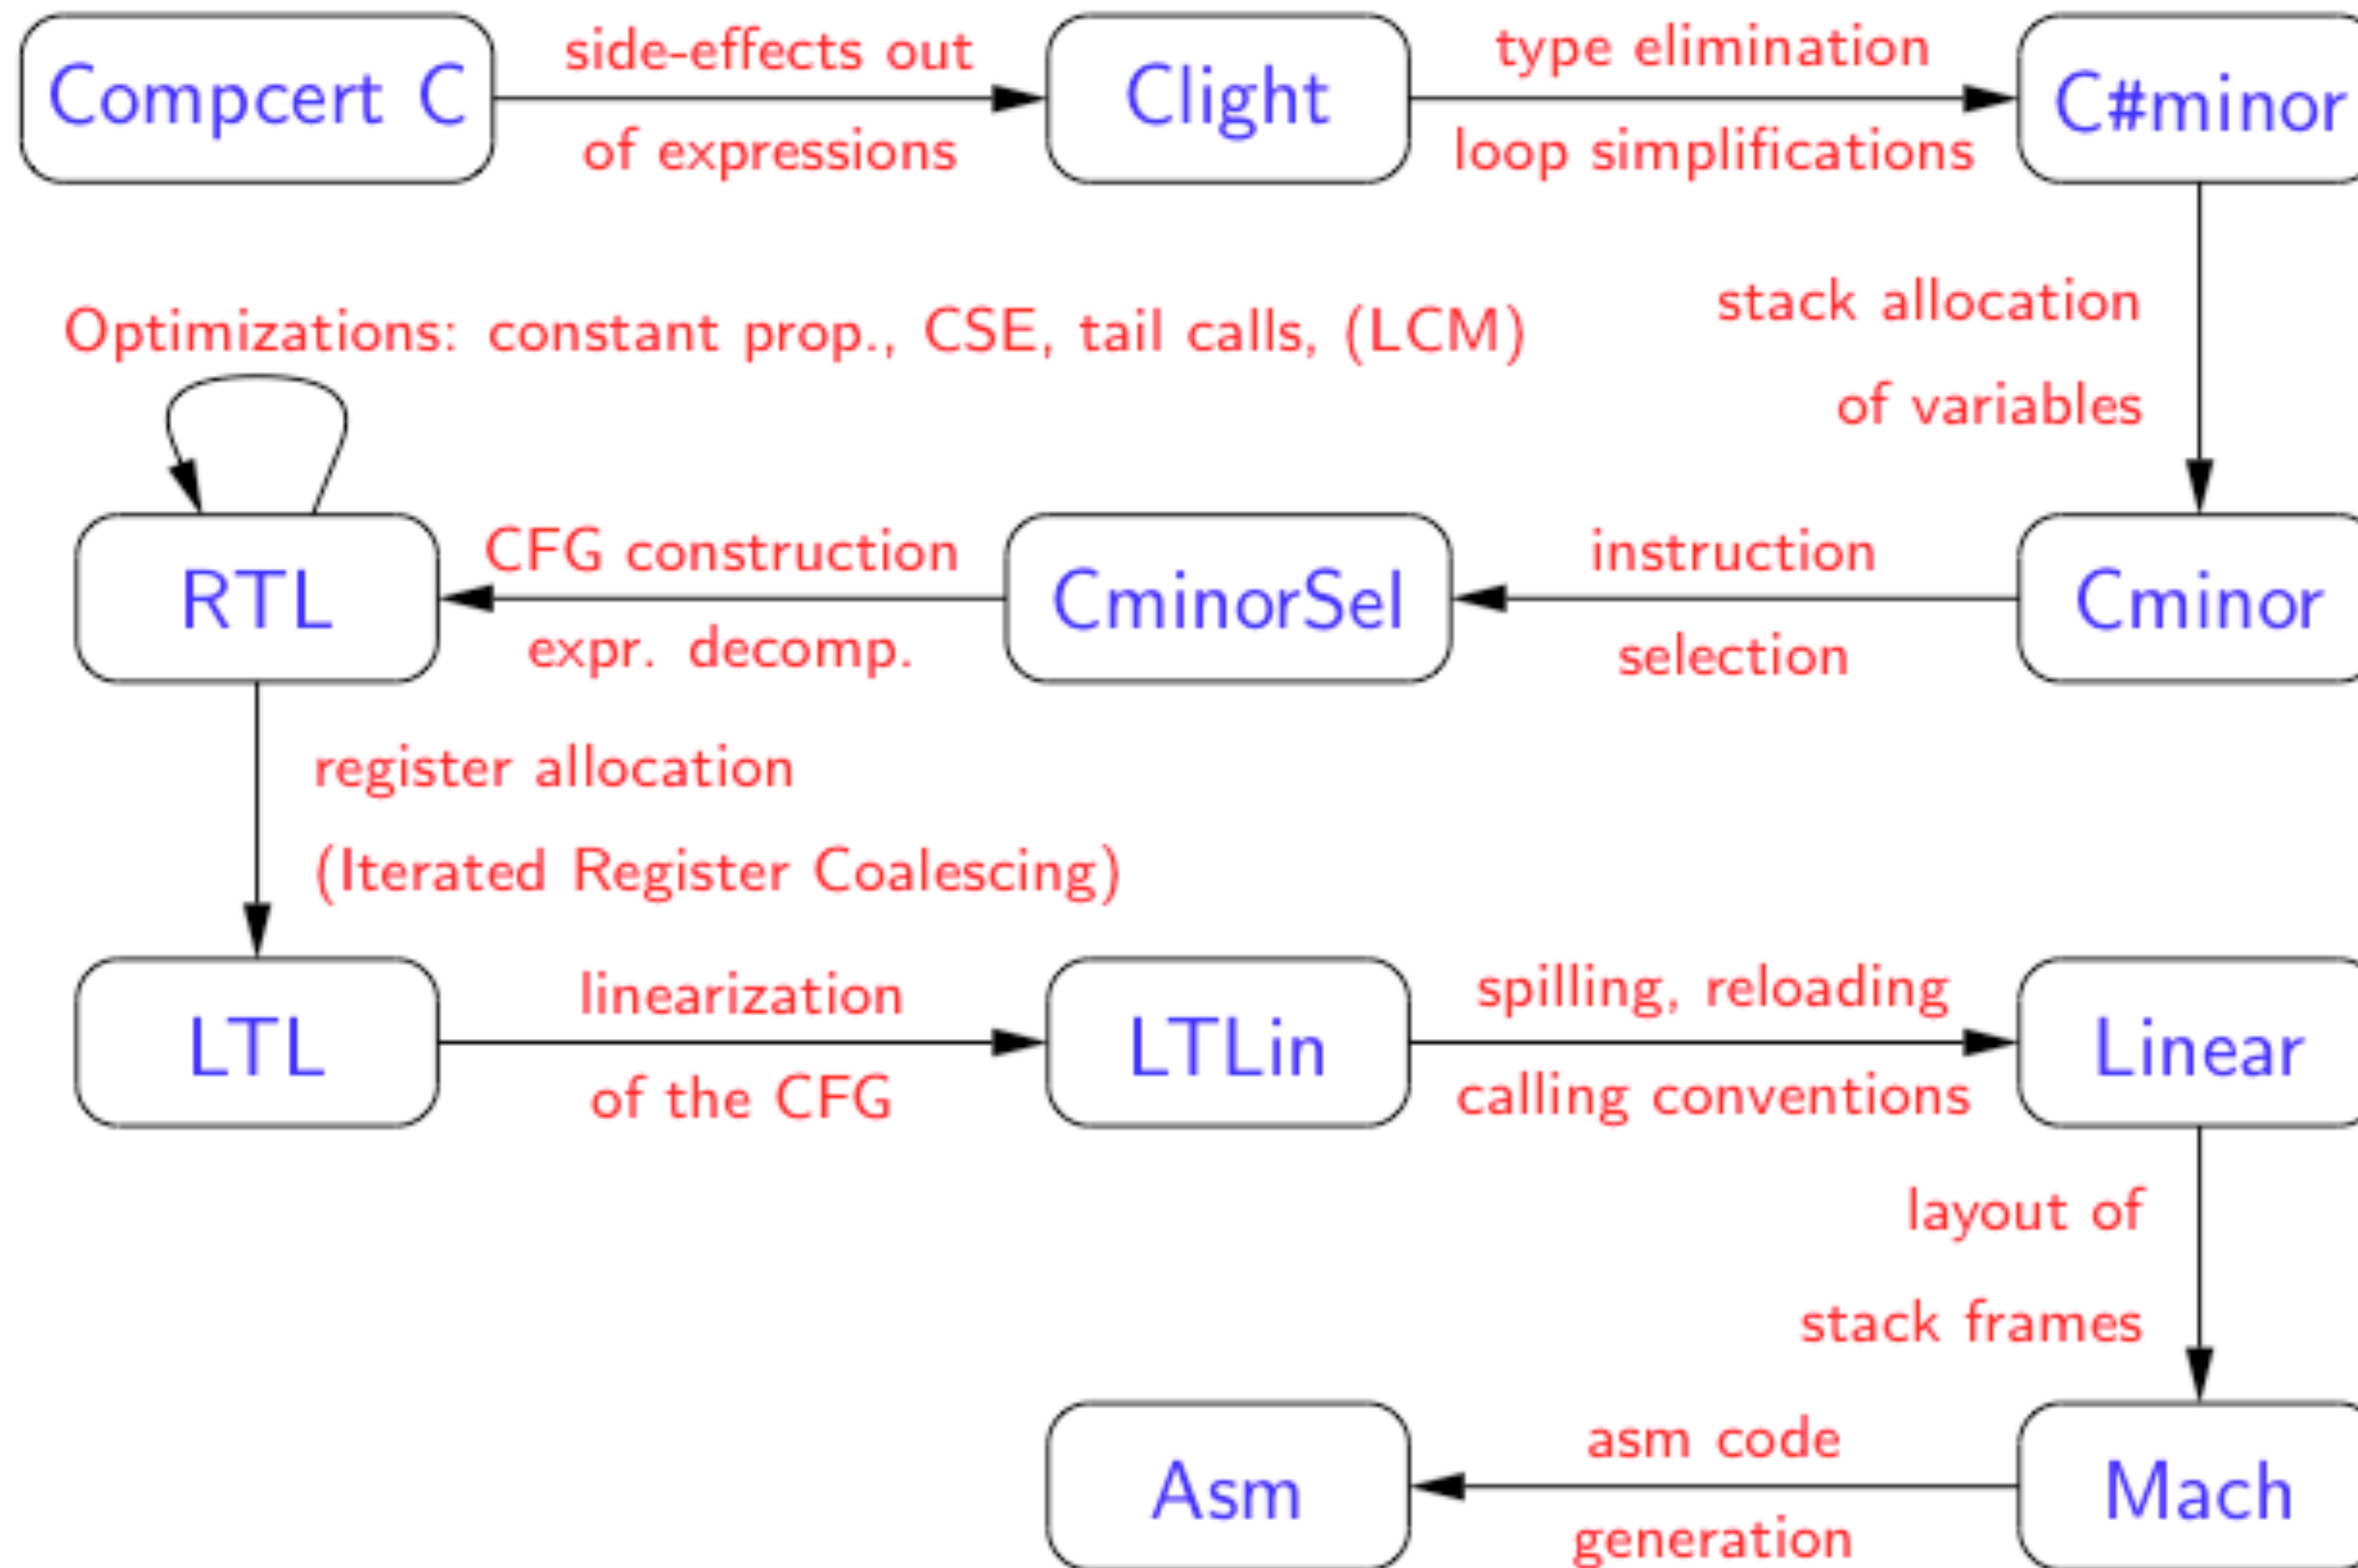
Is this the right semantics?

Testing: for *many* programs p (and inputs i) *test* that

$$\text{run}(p)(i) == v$$

Requires specifying desired $\langle p, i, v \rangle$ combinations
(aka unit testing)

The CompCert C Compiler



Compiler Construction Courses of the Future

Language Specification

syntax definition
name binding
type system
dynamic semantics
translation
transformation

safety properties

Language Implementation

generating implementations
from specifications

parser generation
constraint resolution
partial evaluation

...

Language Verification

proving correctness

Language Testing

test generation

Operational Semantics

Operational Semantics

What is the result of execution of a program?

- How is that result achieved?

Natural Semantics

- How is overall result of execution obtained?

Structural Operational Semantics

- What are the individual steps of an execution?

Defined using a transition system

DynSem

Design and implementation of DynSem, a domain-specific language for the definition of the dynamic semantics of programming languages.

DynSem provides facilities for implicitly passing memory (environment, store).

DynSem specifications can be compiled to or interpreted as interpreters for the languages they define.

RTA 2015

<http://dx.doi.org/10.4230/LIPICs.RTA.2015.365>

DynSem: A DSL for Dynamic Semantics Specification

Vlad Vergu, Pierre Neron, and Eelco Visser

Delft University of Technology
Delft, The Netherlands
{v.a.vergu|p.j.m.neron|visser}@tudelft.nl

Abstract

The formal semantics of a programming language and its implementation are typically separately defined, with the risk of divergence such that properties of the formal semantics are not properties of the implementation. In this paper, we present DynSem, a domain-specific language for the specification of the dynamic semantics of programming languages that aims at supporting both formal reasoning and efficient interpretation. DynSem supports the specification of the *operational semantics* of a language by means of *statically typed conditional term reduction rules*. DynSem supports *concise* specification of reduction rules by providing *implicit build and match coercions* based on reduction arrows and implicit term constructors. DynSem supports *modular* specification by adopting implicit propagation of semantic components from I-MSOS, which allows omitting propagation of components such as environments and stores from rules that do not affect those. DynSem supports the declaration of *native operators* for delegation of aspects of the semantics to an external definition or implementation. DynSem supports the definition of auxiliary *meta-functions*, which can be expressed using regular reduction rules and are subject to semantic component propagation. DynSem specifications are *executable* through automatic generation of a Java-based AST interpreter.

1998 ACM Subject Classification F.3.2. Semantics of Programming Languages (D.3.1.)

Keywords and phrases programming languages, dynamic semantics, reduction semantics, semantics engineering, IDE, interpreters, modularity

Digital Object Identifier 10.4230/LIPICs.RTA.2015.365

1 Introduction

The specification of the dynamic semantics is the core of a programming language design as it describes the runtime behavior of programs. In practice, the implementation of a compiler or an interpreter for the language often stands as the *only* definition of the semantics of a language. Such implementations, in a traditional programming language, often lack the clarity and the conciseness that a specification in a formal semantics framework provides. Therefore, they are a poor source of *documentation* about the semantics. On the other hand, formal definitions are not executable to the point that they can be used as implementations to run programs. Even when both a formal specification and an implementation co-exist, they typically diverge. As a result, important properties of a language as established based on its formal semantics may not hold for its implementation. Our goal is to unify the semantics engineering and language engineering of programming language designs [22] by providing a notation for the specification of the dynamic semantics that can serve both as a readable formalization as well as the source of an execution engine.

In this paper, we present DynSem, a DSL for the concise, modular, statically typed, and executable specification of the dynamic semantics of programming languages. DynSem

Interpreters for Spoofax Languages

gcdAB.pbox	gcdAB.aterm	gcdAB.evaluated.aterm
<pre>1 let 2 gcd = box(0) 3 in 4 let f = 5 fun (a, b) { 6 if (b == 0) 7 a 8 else 9 unbox(gcd)(b, a % b) 10 end 11 } 12 in 13 setbox(gcd, f); 14 unbox(gcd)(1134903170, 1836) 15 end 16 end 17</pre>	<pre>1 Let(2 [Bind("gcd", Box(Num("0")))] 3 , Let(4 [Bind(5 "f" 6 , Fun(7 ["a", "b"] 8 , If(9 Eq(Var("b"), Num("0")) 10 , Var("a") 11 , App(12 Unbox(Var("gcd")) 13 , [Var("b"), Mod(Var("a"), Var("b"))]) 14) 15) 16) 17] 18 , Seq(19 SetBox(Var("gcd"), Var("f")) 20 , App(21 Unbox(Var("gcd")) 22 , [Num("1134903170"), Num("1836")] 23) 24) 25) 26) 27)</pre>	<pre>1 R_default_V(2 NumV(34) 3 , Map(4 "Store" 5 , Bind(923, RefV(925)) 6 , Bind(7 925 8 , ObjV(9 Map(10 "Env" 11 , Bind("outer", 922) 12 , Bind("self", 923) 13 , Bind("super", 924) 14) 15) 16) 17 , Bind(18 926 19 , ClosV(20 Fun(21 ["a", "b"] 22 , If(23 Eq(Var("b"), Num("0")) 24 , Var("a") 25 , App(26 Unbox(Var("gcd")) 27 , [Var("b"), Mod(Var("a"), Var("b"))] 28) 29) 30) 31) 32) 33)</pre>

Example: DynSem Semantics of PAPL-Box

```
let
  fac = box(0)
in
  let f = fun (n) {
    if (n == 0)
      1
    else
      n * (unbox(fac) (n - 1))
    end
  }
  in
    setbox(fac, f);
    unbox(fac)(10)
  end
end
```

Features

- Arithmetic
- Booleans
- Comparisons
- Mutable variables
- Functions
- Boxes

Components

- Syntax in SDF3
- Dynamic Semantics in DynSem

Abstract Syntax from Concrete Syntax

```
module Arithmetic
```

```
imports Expressions
```

```
imports Common
```

```
context-free syntax
```

```
Expr.Num    = INT
```

```
Expr.Plus   = [[Expr] + [Expr]] {left}
```

```
Expr.Minus  = [[Expr] - [Expr]] {left}
```

```
Expr.Times  = [[Expr] * [Expr]] {left}
```

```
Expr.Mod    = [[Expr] % [Expr]] {left}
```

```
context-free priorities
```

```
{left: Expr.Times Expr.Mod }
```

```
> {left: Expr.Minus Expr.Plus }
```

```
module Arithmetic-sig
```

```
imports Expressions-sig
```

```
imports Common-sig
```

```
signature
```

```
sorts Expr
```

```
constructors
```

```
Num    : INT -> Expr
```

```
Plus   : Expr * Expr -> Expr
```

```
Minus  : Expr * Expr -> Expr
```

```
Times  : Expr * Expr -> Expr
```

```
Mod    : Expr * Expr -> Expr
```

src-gen/ds-signatures/Arithmetic-sig

Values, Meta-Variables, and Arrows

```
module values
```

```
signature
```

```
  sorts V Unit
```

```
  constructors
```

```
    U : Unit
```

```
  variables
```

```
    v : V
```

```
module expressions
```

```
imports values
```

```
imports Expressions-sig
```

```
signature
```

```
  arrows
```

```
    Expr --> V
```

```
  variables
```

```
    e : Expr
```

```
    x : String
```

Term Reduction Rules

```
module arithmetic-explicit

imports expressions primitives Arithmetic-sig

signature
  constructors
    NumV: Int -> V

rules

  Num(__String2INT__(n)) --> NumV(str2int(n)).

  Plus(e1, e2) --> NumV(plusI(i1, i2))
  where
    e1 --> NumV(i1); e2 --> NumV(i2).

  Minus(e1, e2) --> NumV(minusI(i1, i2))
  where
    e1 --> NumV(i1); e2 --> NumV(i2).
```

```
module expressions

imports values
imports Expressions-sig

signature
  arrows
    Expr --> V
  variables
    e : Expr
    x : String
```

Native Operations

```
module arithmetic-explicit
imports expressions primitives Arithmetic-sig

signature
  constructors
    NumV: Int -> V

rules

  Num(__String2INT__(n)) --> NumV(str2int(n)).

  Plus(e1, e2) --> NumV(plusI(i1, i2))
  where
    e1 --> NumV(i1); e2 --> NumV(i2).

  Minus(e1, e2) --> NumV(minusI(i1, i2))
  where
    e1 --> NumV(i1); e2 --> NumV(i2).
```

```
module primitives
signature
  native operators
    str2int : String -> Int
    plusI   : Int * Int -> Int
    minusI  : Int * Int -> Int
```

```
public class Natives {

  public static int plusI_2(int i1, int i2) {
    return i1 + i2;
  }

  public static int str2int_1(String s) {
    return Integer.parseInt(s);
  }
}
```


Arrows as Coercions

rules

`Plus(NumV(i1), NumV(i2)) --> NumV(plusI(i1, i2)).`

signature


constructors

`Plus : Expr * Expr -> Expr`

`NumV : Int -> V`

arrows

`Expr --> V`



rules

`Plus(e1, e2) --> NumV(plusI(i1, i2))`

where

`e1 --> NumV(i1);`

`e2 --> NumV(i2).`

Modular

```
module arithmetic

imports Arithmetic-sig
imports expressions
imports primitives
```

```
signature
  constructors
    NumV: Int -> V
```

```
rules

Num(str) --> NumV(str2int(str)).

Plus(NumV(i1), NumV(i2)) --> NumV(plusI(i1, i2)).

Minus(NumV(i1), NumV(i2)) --> NumV(minusI(i1, i2)).

Times(NumV(i1), NumV(i2)) --> NumV(timesI(i1, i2)).

Mod(NumV(i1), NumV(i2)) --> NumV(modI(i1, i2)).
```

```
module boolean

imports Booleans-sig expressions
```

```
signature
  constructors
    BoolV : Bool -> V
```

```
rules

True() --> BoolV(true).
False() --> BoolV(false).

Not(BoolV(false)) --> BoolV(true).
Not(BoolV(true)) --> BoolV(false).

Or(BoolV(true), _) --> BoolV(true).
Or(BoolV(false), e) --> e.

And(BoolV(false), _) --> BoolV(false).
And(BoolV(true), e) --> e.
```

```
module comparison

imports Comparisons-sig arithmetic boolean
```

```
rules

Gt(NumV(i1), NumV(i2)) --> BoolV(gtI(i1, i2)).

Eq(NumV(i1), NumV(i2)) --> BoolV(eqI(i1, i2)).

Eq(BoolV(b1), BoolV(b2)) --> BoolV(eqB(b1, b2)).
```

Interpreter is not defined as a single match over sort

Control-Flow

```
module controlflow


imports ControlFlow-sig
imports expressions
imports boolean

rules

  Seq(v, e2) --> e2.

  If(BoolV(true), e1, _) --> e1.

  If(BoolV(false), _, e2) --> e2.
```



```
module controlflow

imports ControlFlow-sig
imports expressions
imports boolean

rules

  Seq(e1, e2) --> v2
  where
    e1 --> v1;
    e2 --> v2.

  If(e1, e2, e3) --> v
  where
    e1 --> BoolV(true);
    e2 --> v.

  If(e1, e2, e3) --> v
  where
    e1 --> BoolV(false);
    e3 --> v.
```

Immutable Variables: Environment Passing

constructors

```
Let : ID * Expr * Expr -> Expr  
Var : ID -> Expr
```

module variables

```
imports Variables-sig environment
```

rules

```
E |- Let(x, v: V, e2) --> v2  
where  
  Env {x |--> v, E} |- e2 --> v2.  
  
E |- Var(x) --> E[x].
```

module environment

```
imports values
```

signature

```
sort aliases
```

```
Env = Map<String, V>
```

```
variables
```

```
E : Env
```

First-Class Functions: Environment in Closure

constructors

```
Fun : ID * Expr -> Expr  
App : Expr * Expr -> Expr
```

module unary-functions

imports expressions environment

signature

constructors

```
ClosV : String * Expr * Env -> V
```

rules

```
E |- Fun(x, e) --> ClosV(x, e, E).
```

```
E |- App(e1, e2) --> v
```

where

```
E |- e1 --> ClosV(x, e, E');
```

```
E |- e2 --> v2;
```

```
Env {x |--> v2, E'} |- e --> v.
```

module environment

imports values

signature

sort aliases

```
Env = Map<String, V>
```


variables

```
E : Env
```


Implicit Propagation

rules

```
Plus(NumV(i1), NumV(i2)) --> NumV(plusI(i1, i2)).
```




rules

```
Plus(e1, e2) --> NumV(plusI(i1, i2))
```

where

```
e1 --> NumV(i1);
```

```
e2 --> NumV(i2).
```



rules

```
E |- Plus(e1, e2) --> NumV(plusI(i1, i2))
```

where

```
E |- e1 --> NumV(i1);
```

```
E |- e2 --> NumV(i2).
```

Mutable Boxes: Store

```
module box
imports store arithmetic
signature
  constructors
    Box      : Expr -> Expr
    Unbox    : Expr -> Expr
    SetBox   : Expr * Expr -> Expr
  constructors
    BoxV: Int -> V
rules

Box(e) :: S --> BoxV(loc) :: Store {loc |--> v, S'}
where e :: S --> v :: S';
      fresh => loc.

Unbox(BoxV(loc)) :: S --> S[loc].

SetBox(BoxV(loc), v) :: S --> v :: Store {loc |--> v, S}.
```

```
module store
imports values

signature
  sort aliases
    Store = Map<Int, V>
  variables
    S : Store
```

Mutable Variables: Environment + Store

constructors

```
Let : ID * Expr * Expr -> Expr
Var : ID -> Expr
Set : String * Expr -> Expr
```

```
module variables-mutable
imports Variables-sig store
rules
```

```
E |- Var(x) :: S --> v :: S
where E[x] => loc; S[loc] => v.
```

```
E |- Let(x, v, e2) :: S1 --> v2 :: S3
where
  fresh => loc;
  {loc |--> v, S1} => S2;
  Env {x |--> loc, E} |- e2 :: S2 --> v2 :: S3.
```

```
E |- Set(x, v) :: S --> v :: Store {loc |--> v, S}
where E[x] => loc.
```

```
module store
```

```
imports values
```

```
signature
```

```
sort aliases
```

```
Env = Map<ID, Int>
```

```
Store = Map<Int, V>
```

```
variables
```


```
E : Env
```

```
S : Store
```

Implicit Store Threading

rules

```
Plus(NumV(i1), NumV(i2)) --> NumV(plusI(i1, i2)).
```




rules

```
Plus(e1, e2) --> NumV(plusI(i1, i2))
```

where

```
e1 --> NumV(i1);
```

```
e2 --> NumV(i2).
```



rules

```
E |- Plus(e1, e2) :: S1 --> NumV(plusI(i1, i2)) :: S3
```

where

```
E |- e1 :: S1 --> NumV(i1) :: S2;
```

```
E |- e2 :: S2 --> NumV(i2) :: S3.
```

Abstraction: Env/Store Meta Functions

```
module store

imports values

signature
  sort aliases
    Env = Map<String, Int>
    Store = Map<Int, V>
  variables
    E : Env
    S : Store
  arrows
    readVar  : String --> V
    bindVar  : String * V --> Env
    writeVar : String * V --> V

    allocate : V --> Int
    write    : Int * V --> V
    read     : Int --> V
```

```
rules

allocate(v) --> loc
where
  fresh ==> loc;
  write(loc, v) --> ..

write(loc, v) :: S -->
  v :: Store {loc |--> v, S}.

read(loc) :: S --> S[loc] :: S.

rules

bindVar(x, v) --> {x |--> loc}
where allocate(v) --> loc.

E |- readVar(x) --> read(E[x]).

E |- writeVar(x, v) --> write(E[x], v).
```

Boxes with Env/Store Meta Functions

```
module boxes

signature
  constructors
    Box      : Expr -> Expr
    Unbox    : Expr -> Expr
    SetBox   : Expr * Expr -> Expr
  constructors
    BoxV: V -> V

rules

  Box(v) --> BoxV(NumV(allocate(v))).

  Unbox(BoxV(NumV(loc))) --> read(loc).

  SetBox(BoxV(NumV(loc)), v) --> write(loc, v).
```

Mutable Variables with Env/Store Meta Functions

constructors

```
Let  : String * Expr * Expr -> Expr
Var  : String -> Expr
Set  : String * Expr -> Expr
```

module variables

imports expressions store

rules

```
Var(x) --> readVar(x).
```

```
E |- Let(x, v1, e) --> v2
```

where

```
  bindVar(x, v1) --> E';
  Env {E', E} |- e --> v2.
```

```
Set(x, v) --> v
```

where

```
  writeVar(x, v) --> _.
```


Functions with Multiple Arguments

```
module functions

imports Functions-sig
imports variables

signature
  constructors
    ClosV      : List(ID) * Expr * Env -> V
    bindArgs   : List(ID) * List(Expr) --> Env

rules

  E |- Fun(xs, e) --> ClosV(xs, e, E).

  App(ClosV(xs, e_body, E_clos), es) --> v'
  where
    bindArgs(xs, es) --> E_params;
    Env {E_params, E_clos} |- e_body --> v'.

  bindArgs([], []) --> {}.

  bindArgs([x | xs], [e | es]) --> {E, E'}
  where
    bindVar(x, e) --> E;
    bindArgs(xs, es) --> E'.
```

Tiger in DynSem

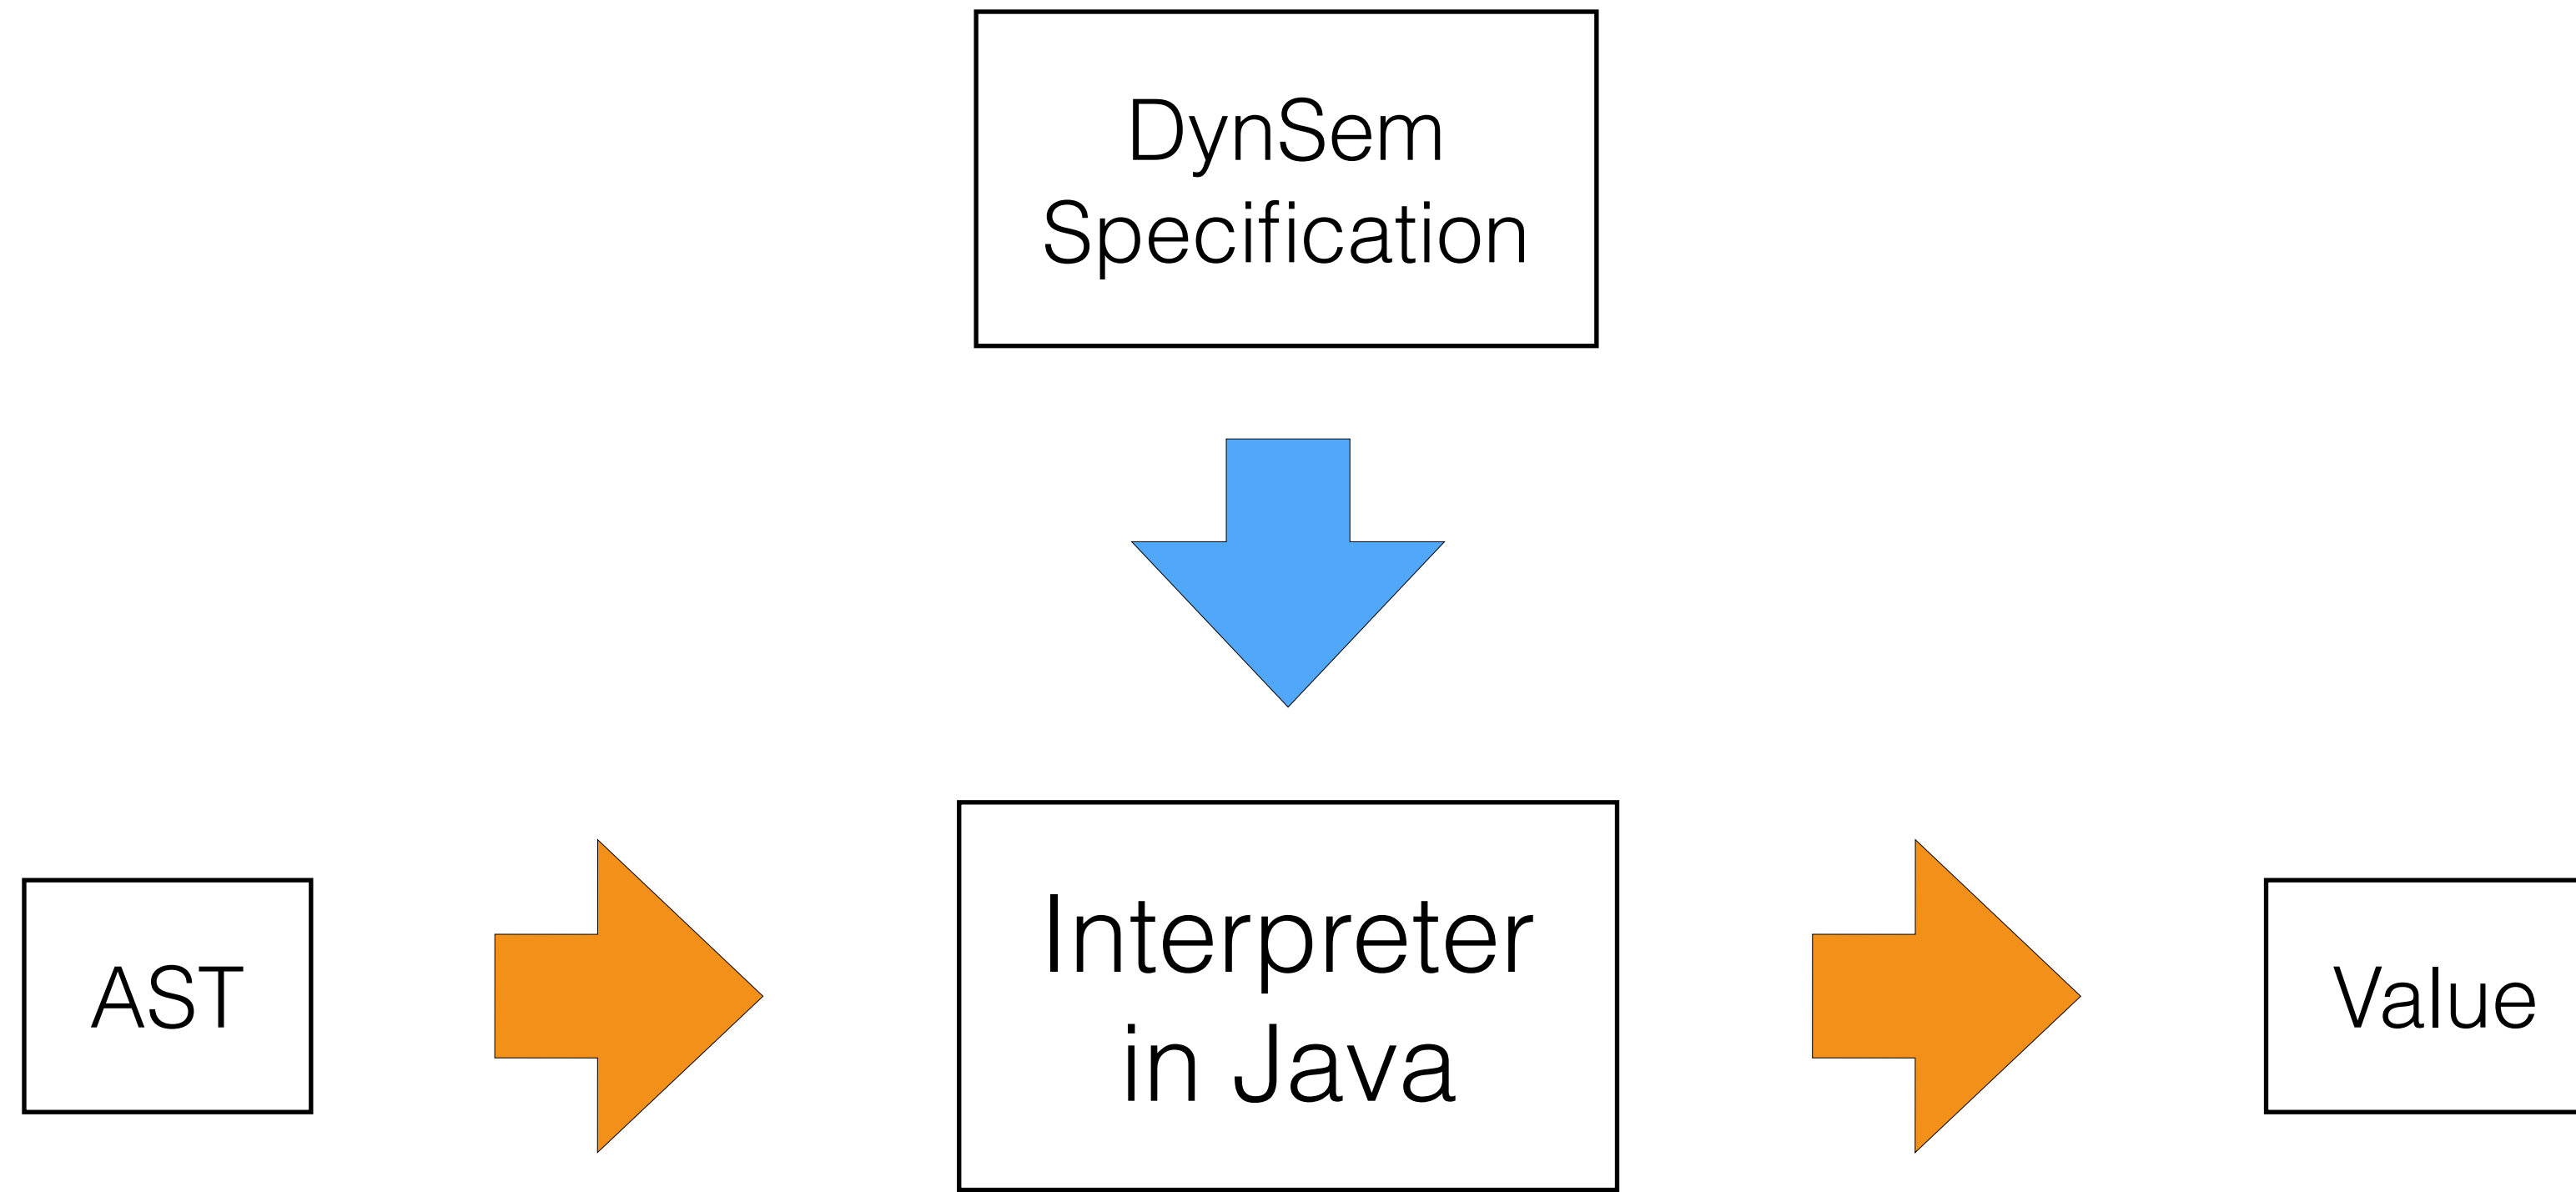
```

store.ds
2
3 imports dynamics/values
4
5 signature // lvalue
6 sorts LValue
7 arrows
8   LValue -lval-> Int
9 variables
10   lv : LValue
11
12 signature // environment
13 sorts Id
14 sort aliases
15   // Address = Int
16   Env = Map(Id, Int)
17 variables
18   a : Int
19 components
20   E : Env
21 arrows
22   lookup(Id) --> Int
23   bind(Id, Int) --> Env
24
25
26 rules
27
28 E |- lookup(x) --> E[x].
29
30 E |- bind(x, a) --> {x l--> a, E}.
31
32 signature // heap
33 sort aliases
34   Heap = Map(Int, V)
35 components
36   H : Heap
37 arrows
38   read(Int) --> V
39   allocate(V) --> Int
40   write(Int, V) --> V
41
42 rules
43
44 read(a) :: H --> H[a].
45
46 write(a, v) :: H --> v :: H {a l--> v, H}
47
48 allocate(v) --> a
49 where
50   fresh => a;
51   write(a, v) --> ..
52
numbers.ds
1 module functions
2
3 imports ds-signatures/Functions-sig
4 imports dynamics/base
5 imports dynamics/store
6 imports dynamics/bindings
7
8 signature
9 constructors
10   ClosureV : List(FArg) * Exp * Env -> V
11 arrows
12   E |- funEnv(List(FunDec)) :: H --> Env :: H
13   E |- evalFuns(List(FunDec)) :: H --> Env :: H
14   E |- evalArgs(List(FArg), List(Exp)) :: H --> Env ::
15
16 rules // function definition
17
18 FunDecs(fds) --> E
19 where
20   funEnv(fds) --> E;
21   E |- evalFuns(fds) --> ..
22
23 E |- funEnv([]) --> E.
24
25 funEnv([FunDec(f, _, _, _) | fds]) --> E
26 where
27   E bindVar(f, UndefV()) |- funEnv(fds) --> E.
28
29 E |- evalFuns([]) --> E.
30
31 E |- evalFuns([FunDec(f, args, _, e) | fds]) --> evalFu
32 where
33   writeVar(f, ClosureV(args, e, E)) --> ..
34
35 rules // function call
36
37 Call(f, es) --> v
38 where
39   readVar(f) --> ClosureV(args, e, E);
40   evalArgs(args, es) --> E';
41   E {E', E} |- e --> v.
42
43 evalArgs([], []) --> {}.
44
45 evalArgs([FArg(x, _) | args], [x | es]) --> {x l--> a,
46 where
47   allocate(v) --> a;
48   evalArgs(args, es) --> E.
49
50 rules // procedure definition
51
functions.ds
1 module functions
2
3 imports ds-signatures/Functions-sig
4 imports dynamics/base
5 imports dynamics/store
6 imports dynamics/bindings
7
8 signature
9 constructors
10   ClosureV : List(FArg) * Exp * Env -> V
11 arrows
12   E |- funEnv(List(FunDec)) :: H --> Env :: H
13   E |- evalFuns(List(FunDec)) :: H --> Env :: H
14   E |- evalArgs(List(FArg), List(Exp)) :: H --> Env ::
15
16 rules // function definition
17
18 FunDecs(fds) --> E
19 where
20   funEnv(fds) --> E;
21   E |- evalFuns(fds) --> ..
22
23 E |- funEnv([]) --> E.
24
25 funEnv([FunDec(f, _, _, _) | fds]) --> E
26 where
27   E bindVar(f, UndefV()) |- funEnv(fds) --> E.
28
29 E |- evalFuns([]) --> E.
30
31 E |- evalFuns([FunDec(f, args, _, e) | fds]) --> evalFu
32 where
33   writeVar(f, ClosureV(args, e, E)) --> ..
34
35 rules // function call
36
37 Call(f, es) --> v
38 where
39   readVar(f) --> ClosureV(args, e, E);
40   evalArgs(args, es) --> E';
41   E {E', E} |- e --> v.
42
43 evalArgs([], []) --> {}.
44
45 evalArgs([FArg(x, _) | args], [x | es]) --> {x l--> a,
46 where
47   allocate(v) --> a;
48   evalArgs(args, es) --> E.
49
50 rules // procedure definition
51
equality.ds
1 module functions
2
3 imports ds-signatures/Functions-sig
4 imports dynamics/base
5 imports dynamics/store
6 imports dynamics/bindings
7
8 signature
9 constructors
10   ClosureV : List(FArg) * Exp * Env -> V
11 arrows
12   E |- funEnv(List(FunDec)) :: H --> Env :: H
13   E |- evalFuns(List(FunDec)) :: H --> Env :: H
14   E |- evalArgs(List(FArg), List(Exp)) :: H --> Env ::
15
16 rules // function definition
17
18 FunDecs(fds) --> E
19 where
20   funEnv(fds) --> E;
21   E |- evalFuns(fds) --> ..
22
23 E |- funEnv([]) --> E.
24
25 funEnv([FunDec(f, _, _, _) | fds]) --> E
26 where
27   E bindVar(f, UndefV()) |- funEnv(fds) --> E.
28
29 E |- evalFuns([]) --> E.
30
31 E |- evalFuns([FunDec(f, args, _, e) | fds]) --> evalFu
32 where
33   writeVar(f, ClosureV(args, e, E)) --> ..
34
35 rules // function call
36
37 Call(f, es) --> v
38 where
39   readVar(f) --> ClosureV(args, e, E);
40   evalArgs(args, es) --> E';
41   E {E', E} |- e --> v.
42
43 evalArgs([], []) --> {}.
44
45 evalArgs([FArg(x, _) | args], [x | es]) --> {x l--> a,
46 where
47   allocate(v) --> a;
48   evalArgs(args, es) --> E.
49
50 rules // procedure definition
51
control-flow.ds
1 signature
2 sort aliases
3   Idx = Map(Int, Int)
4 variables
5   I : Idx
6 constructors
7   ArrayV : Idx -> V
8 arrows
9   initArray(Int, Int, V, Idx) --> Idx
10
11 rules
12
13 Array(_, IntV(i), v) --> ArrayV(I)
14 where
15   initArray(0, i, v, {}) --> I.
16
17 initArray(i, j, v, I) --> I'
18 where
19   case ltI(i, j) of {
20     1 =>
21       allocate(v) --> a;
22       initArray(addI(i, 1), j, v, {i l--> a, I})
23     0 =>
24       I => I'
25   }.
26
27 Subscript(a, IntV(i)) -lval-> I[i]
28 where
29   read(a) --> ArrayV(I).
30
arrays.ds
1 signature
2 sort aliases
3   Idx = Map(Int, Int)
4 variables
5   I : Idx
6 constructors
7   ArrayV : Idx -> V
8 arrows
9   initArray(Int, Int, V, Idx) --> Idx
10
11 rules
12
13 Array(_, IntV(i), v) --> ArrayV(I)
14 where
15   initArray(0, i, v, {}) --> I.
16
17 initArray(i, j, v, I) --> I'
18 where
19   case ltI(i, j) of {
20     1 =>
21       allocate(v) --> a;
22       initArray(addI(i, 1), j, v, {i l--> a, I})
23     0 =>
24       I => I'
25   }.
26
27 Subscript(a, IntV(i)) -lval-> I[i]
28 where
29   read(a) --> ArrayV(I).
30
prettyprint.tig
1 let
2
3 type tree = {key: string, left : tree, right: tree}
4
5 function prettyprint(tree: tree) : string =
6 let
7
8   var output := ""
9
10 function write(s: string) =
11   output := concat(output, s)
12
13 function show(n: int, t: tree) =
14 let function indent(s: string) =
15   (write("\n");
16   for i := 1 to n
17   do write(" ");
18   output := concat(output, s))
19 in if t = nil then indent(".")
20 else (indent(t.key);
21       show(n+1, t.left);
22       show(n+1, t.right))
23 end
24
25 in show(0, tree);
prettyprint.aterm
1 Mod(
2   Let(
3     [ TypeDecs(
4       [ TypeDec(
5         "tree"
6         , RecordTy(
7           [ Field("key", Tid("string"))
8           , Field("left", Tid("tree"))
9           , Field("right", Tid("tree"))
10         ]
11       )
12     ]
13   )
14   , FunDecs(
15     [ FunDec(
16       "prettyprint"
17       , [FArg("tree", Tid("tree"))]
18       , Tid("string")
19       , Let(
20         [ VarDecNoType("output", String("\n"))
21         , FunDecs(
22           [ ProcDec(
23             "write"
24             , [FArg("s", Tid("string"))]
25

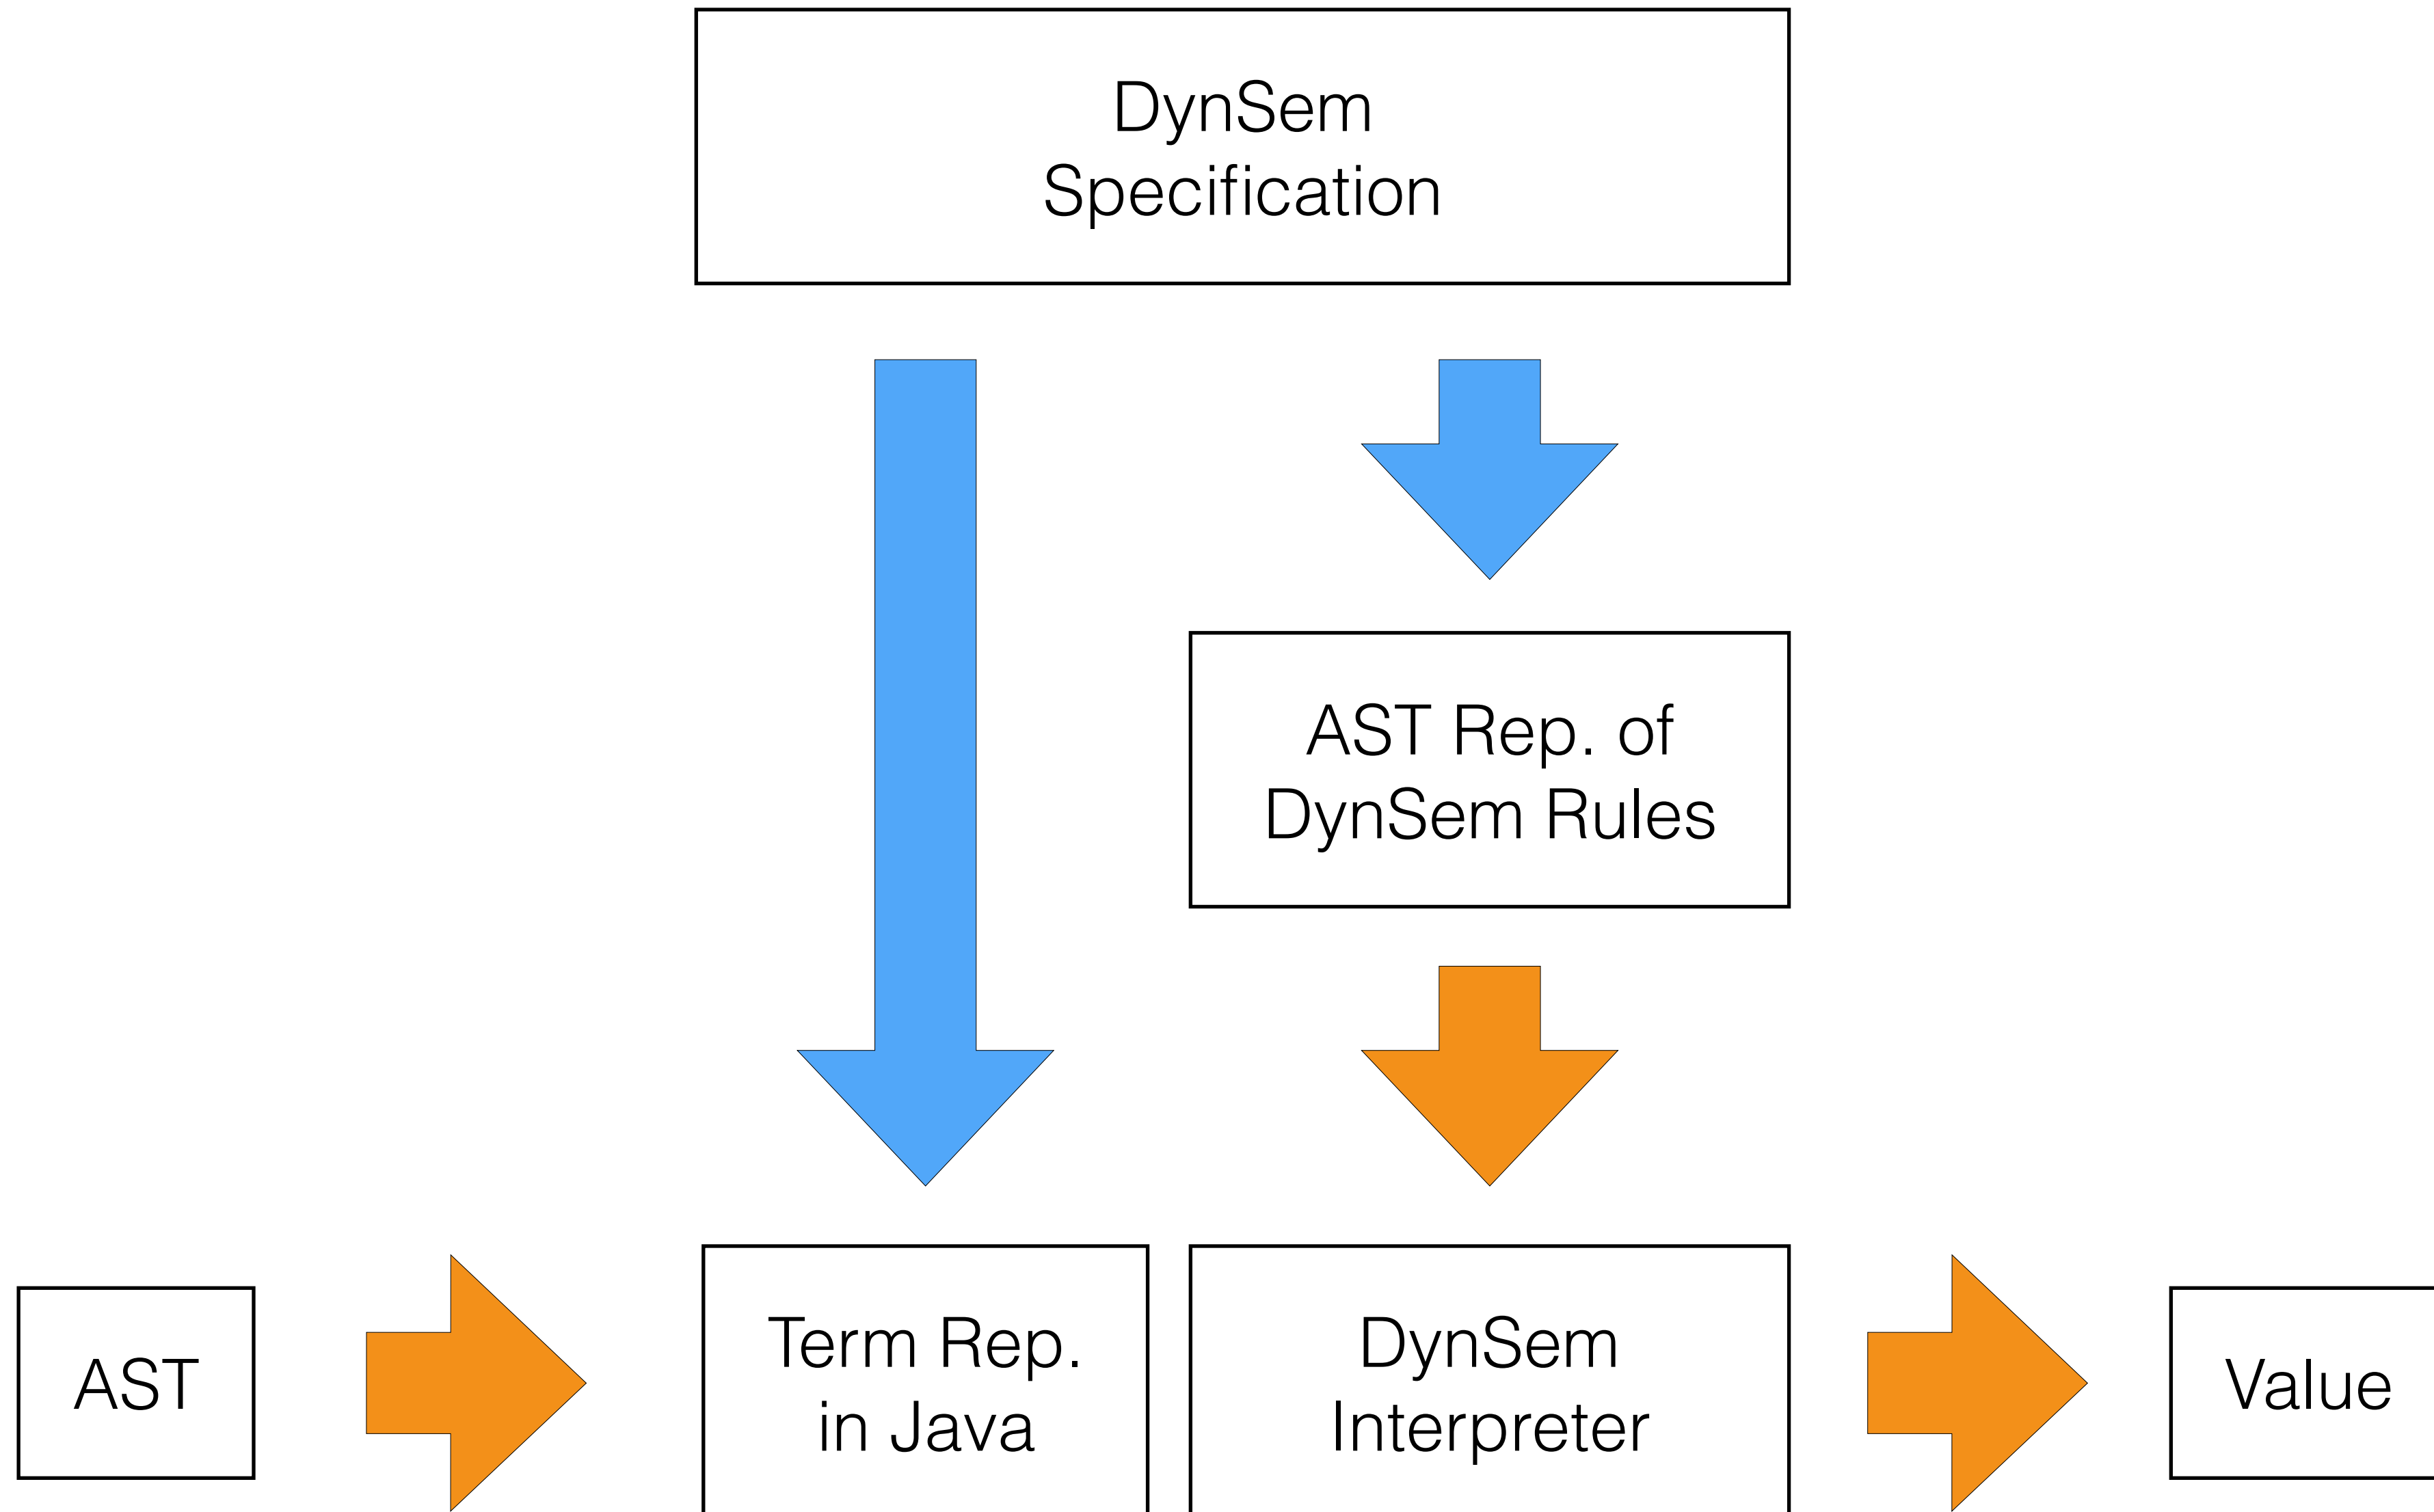
```

Interpreter Generation

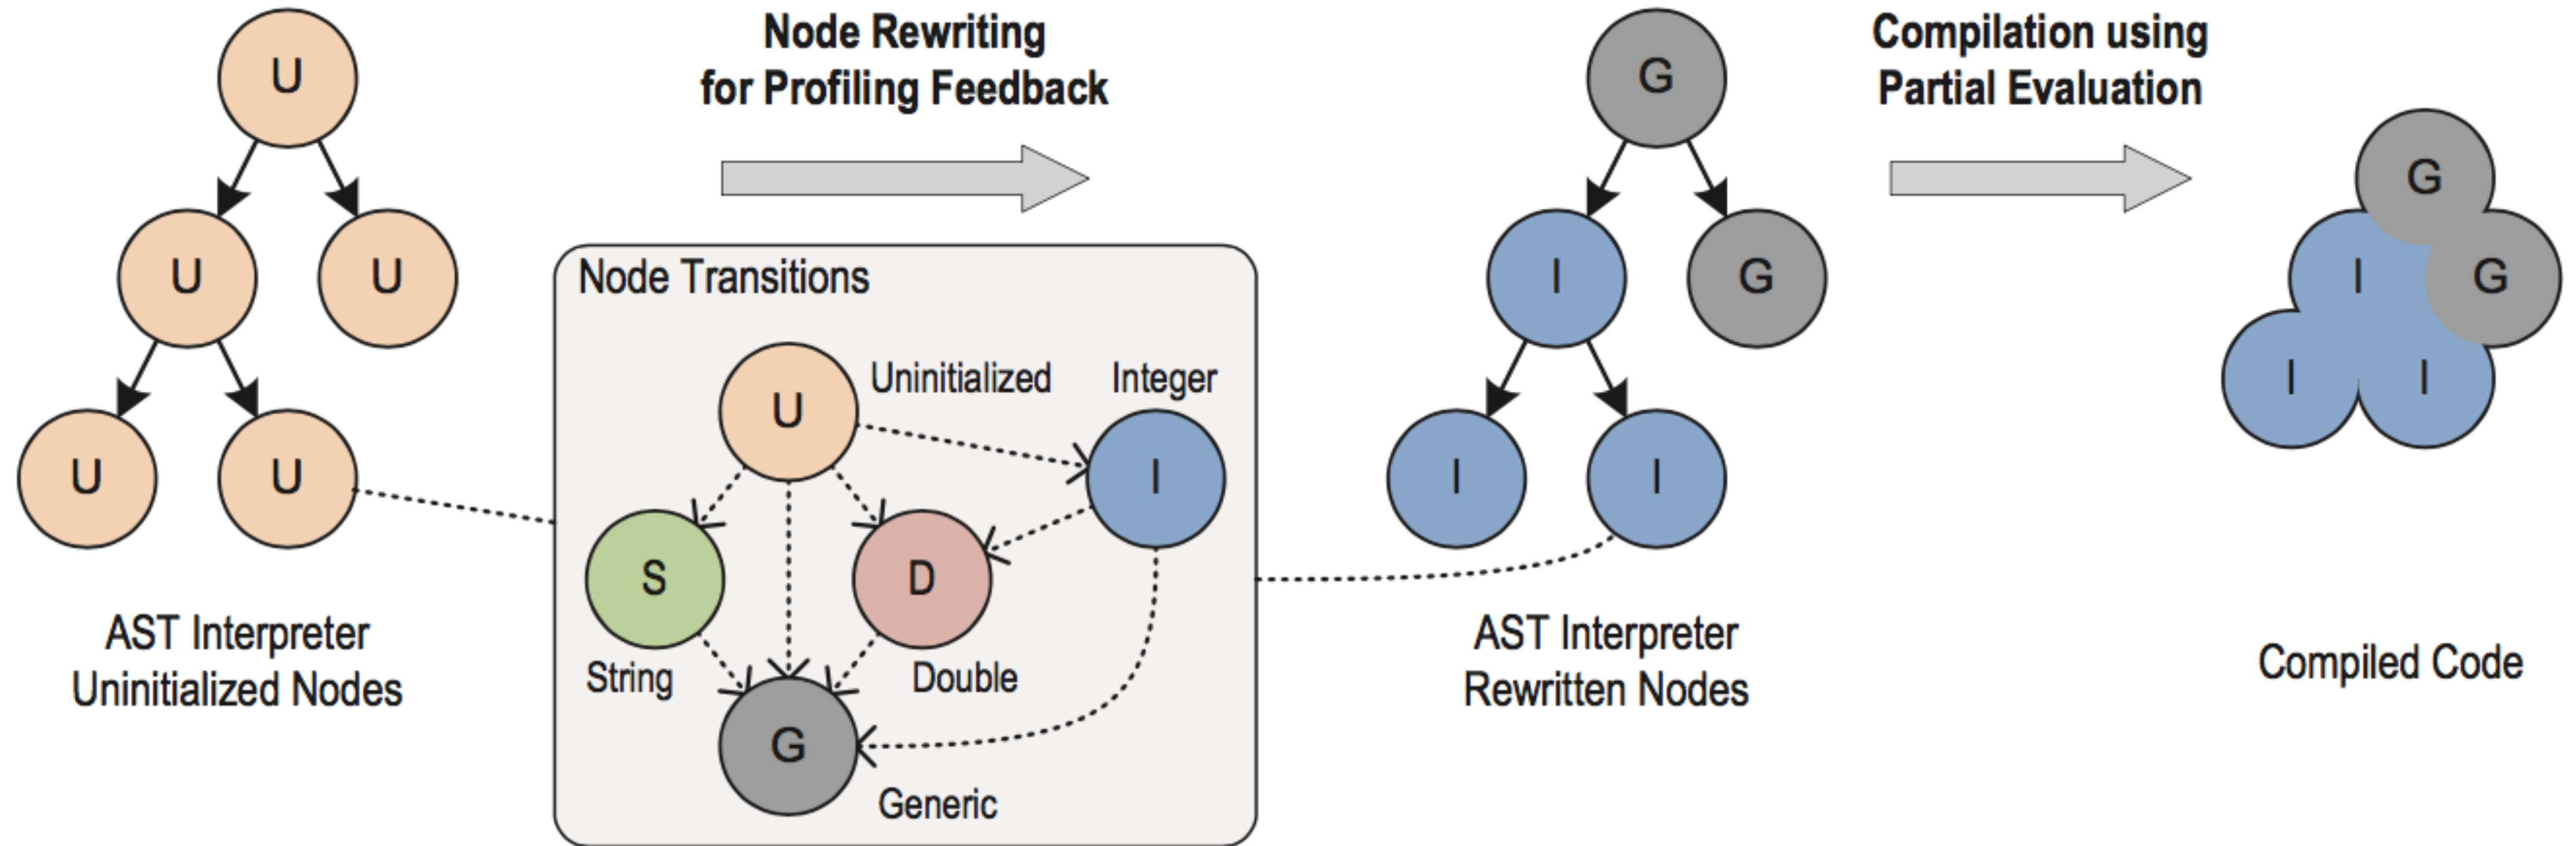
Generating AST Interpreter from Dynamic Semantics



DynSem Meta-Interpreter



Truffle: Partial Evaluation of AST Interpreters



This paper describes the partial evaluation of the DynSem meta-interpreter with an object program.

The implementation makes use of Truffle a framework for run-time (online) partial evaluation.

Truffle gets most of it benefit from Graal, an implementation of the JVM with special support for partial evaluation.

ManLang 2018

<https://doi.org/10.1145/3237009.3237018>

Specializing a Meta-Interpreter

JIT Compilation of DynSem Specifications on the Graal VM

Vlad Vergu
TU Delft
The Netherlands
v.a.vergu@tudelft.nl

Eelco Visser
TU Delft
The Netherlands
visser@acm.org

ABSTRACT

DynSem is a domain-specific language for concise specification of the dynamic semantics of programming languages, aimed at rapid experimentation and evolution of language designs. DynSem specifications can be executed to interpret programs in the language under development. To enable fast turnaround during language development, we have developed a meta-interpreter for DynSem specifications, which requires minimal processing of the specification. In addition to fast development time, we also aim to achieve fast run times for interpreted programs.

In this paper we present the design of a meta-interpreter for DynSem and report on experiments with JIT compiling the application of the meta-interpreter on the Graal VM. By interpreting specifications directly, we have minimal compilation overhead. By specializing pattern matches, maintaining call-site dispatch chains and using native control-flow constructs we gain significant run-time performance. We evaluate the performance of the meta-interpreter when applied to the Tiger language specification running a set of common benchmark programs. Specialization enables the Graal VM to JIT compile the meta-interpreter giving speedups of up to factor 15 over running on the standard Oracle Java VM.

CCS CONCEPTS

• **Software and its engineering** → **Interpreters; Domain specific languages; Semantics;**

KEYWORDS

dynamic semantics, interpretation, JIT, run-time optimization

ACM Reference Format:

Vlad Vergu and Eelco Visser. 2018. Specializing a Meta-Interpreter: JIT Compilation of DynSem Specifications on the Graal VM. In *15th International Conference on Managed Languages & Runtimes (ManLang’18)*, September 12–14, 2018, Linz, Austria. ACM, New York, NY, USA, 14 pages. <https://doi.org/10.1145/3237009.3237018>

1 INTRODUCTION

The dynamic semantics of a programming language defines the run time execution behavior of programs in the language. Ideally,

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.
ManLang’18, September 12–14, 2018, Linz, Austria
© 2018 Association for Computing Machinery.
ACM ISBN 978-1-4503-6424-9/18/09... \$15.00
<https://doi.org/10.1145/3237009.3237018>

the design of a programming language starts with the specification of its dynamic semantics to provide a high-level readable and unambiguous definition. However, understanding the design of a programming language also requires experimentation by actually running programs. Therefore, this ideal route is rarely taken, but language designs are embodied in the implementation of interpreters or compilers instead.

We have previously designed DynSem [33], a high-level meta-DSL for dynamic semantics specifications of programming languages, with the aim of supporting readable *and* executable specification. It supports the definition of modular and concise semantics by means of reduction rules with implicit propagation of contextual information. DynSem’s executable semantics entails that specifications can be used to interpret object language programs.

In our early prototypes, DynSem specifications were compiled to an interpreter. The process of generating a Java implementation of an interpreter and compiling that generated code caused long turnaround times during language prototyping. In order to support rapid prototyping with short turnaround times, we turned to interpreting specifications directly instead of compiling them. A DynSem interpreter is a *meta-interpreter* since the programs it interprets are themselves interpreters. Figure 1 depicts the high-level architecture of the DynSem meta-interpreter. First, a DynSem specification is desugared (explicated) to make implicit passing of semantic components explicit. The resulting specification in DynSem Core is then loaded into the meta-interpreter together with the AST of the interpreted object program. The interpreter consumes the program as input enacting the specification. This produces the desired result of a short turnaround time for experimenting with dynamic semantics specifications.

Meta-interpretation reduces the turnaround time at the expense of execution performance. At run time there are two interpreter layers operating (the meta-language interpreter and the object-language interpreter) which introduces substantial overhead. While we envision DynSem as a convenient way to prototype the dynamic semantics of programming languages, ultimately we also envision it as a convenient way to bridge the gap between the prototyping and production phases of a programming language’s lifecycle. Thus, we not only want an interpreter fast, but we also want a fast interpreter, which raises the question: Can we achieve fast object-language interpreters by optimizing the meta-interpretation of dynamic semantics specifications?

Direct vanilla interpreters are in general slow to begin with, even when they are implemented in a host language that is JIT-ed. This is because the host JIT is unable to see patterns in the object language and to meaningfully optimize the interpreter. The task of optimizing an interpreter has traditionally been long and

Partial Evaluation

Partial Evaluation = Program Specialization

Partial evaluation
=
Specializing a program to its static (known) inputs

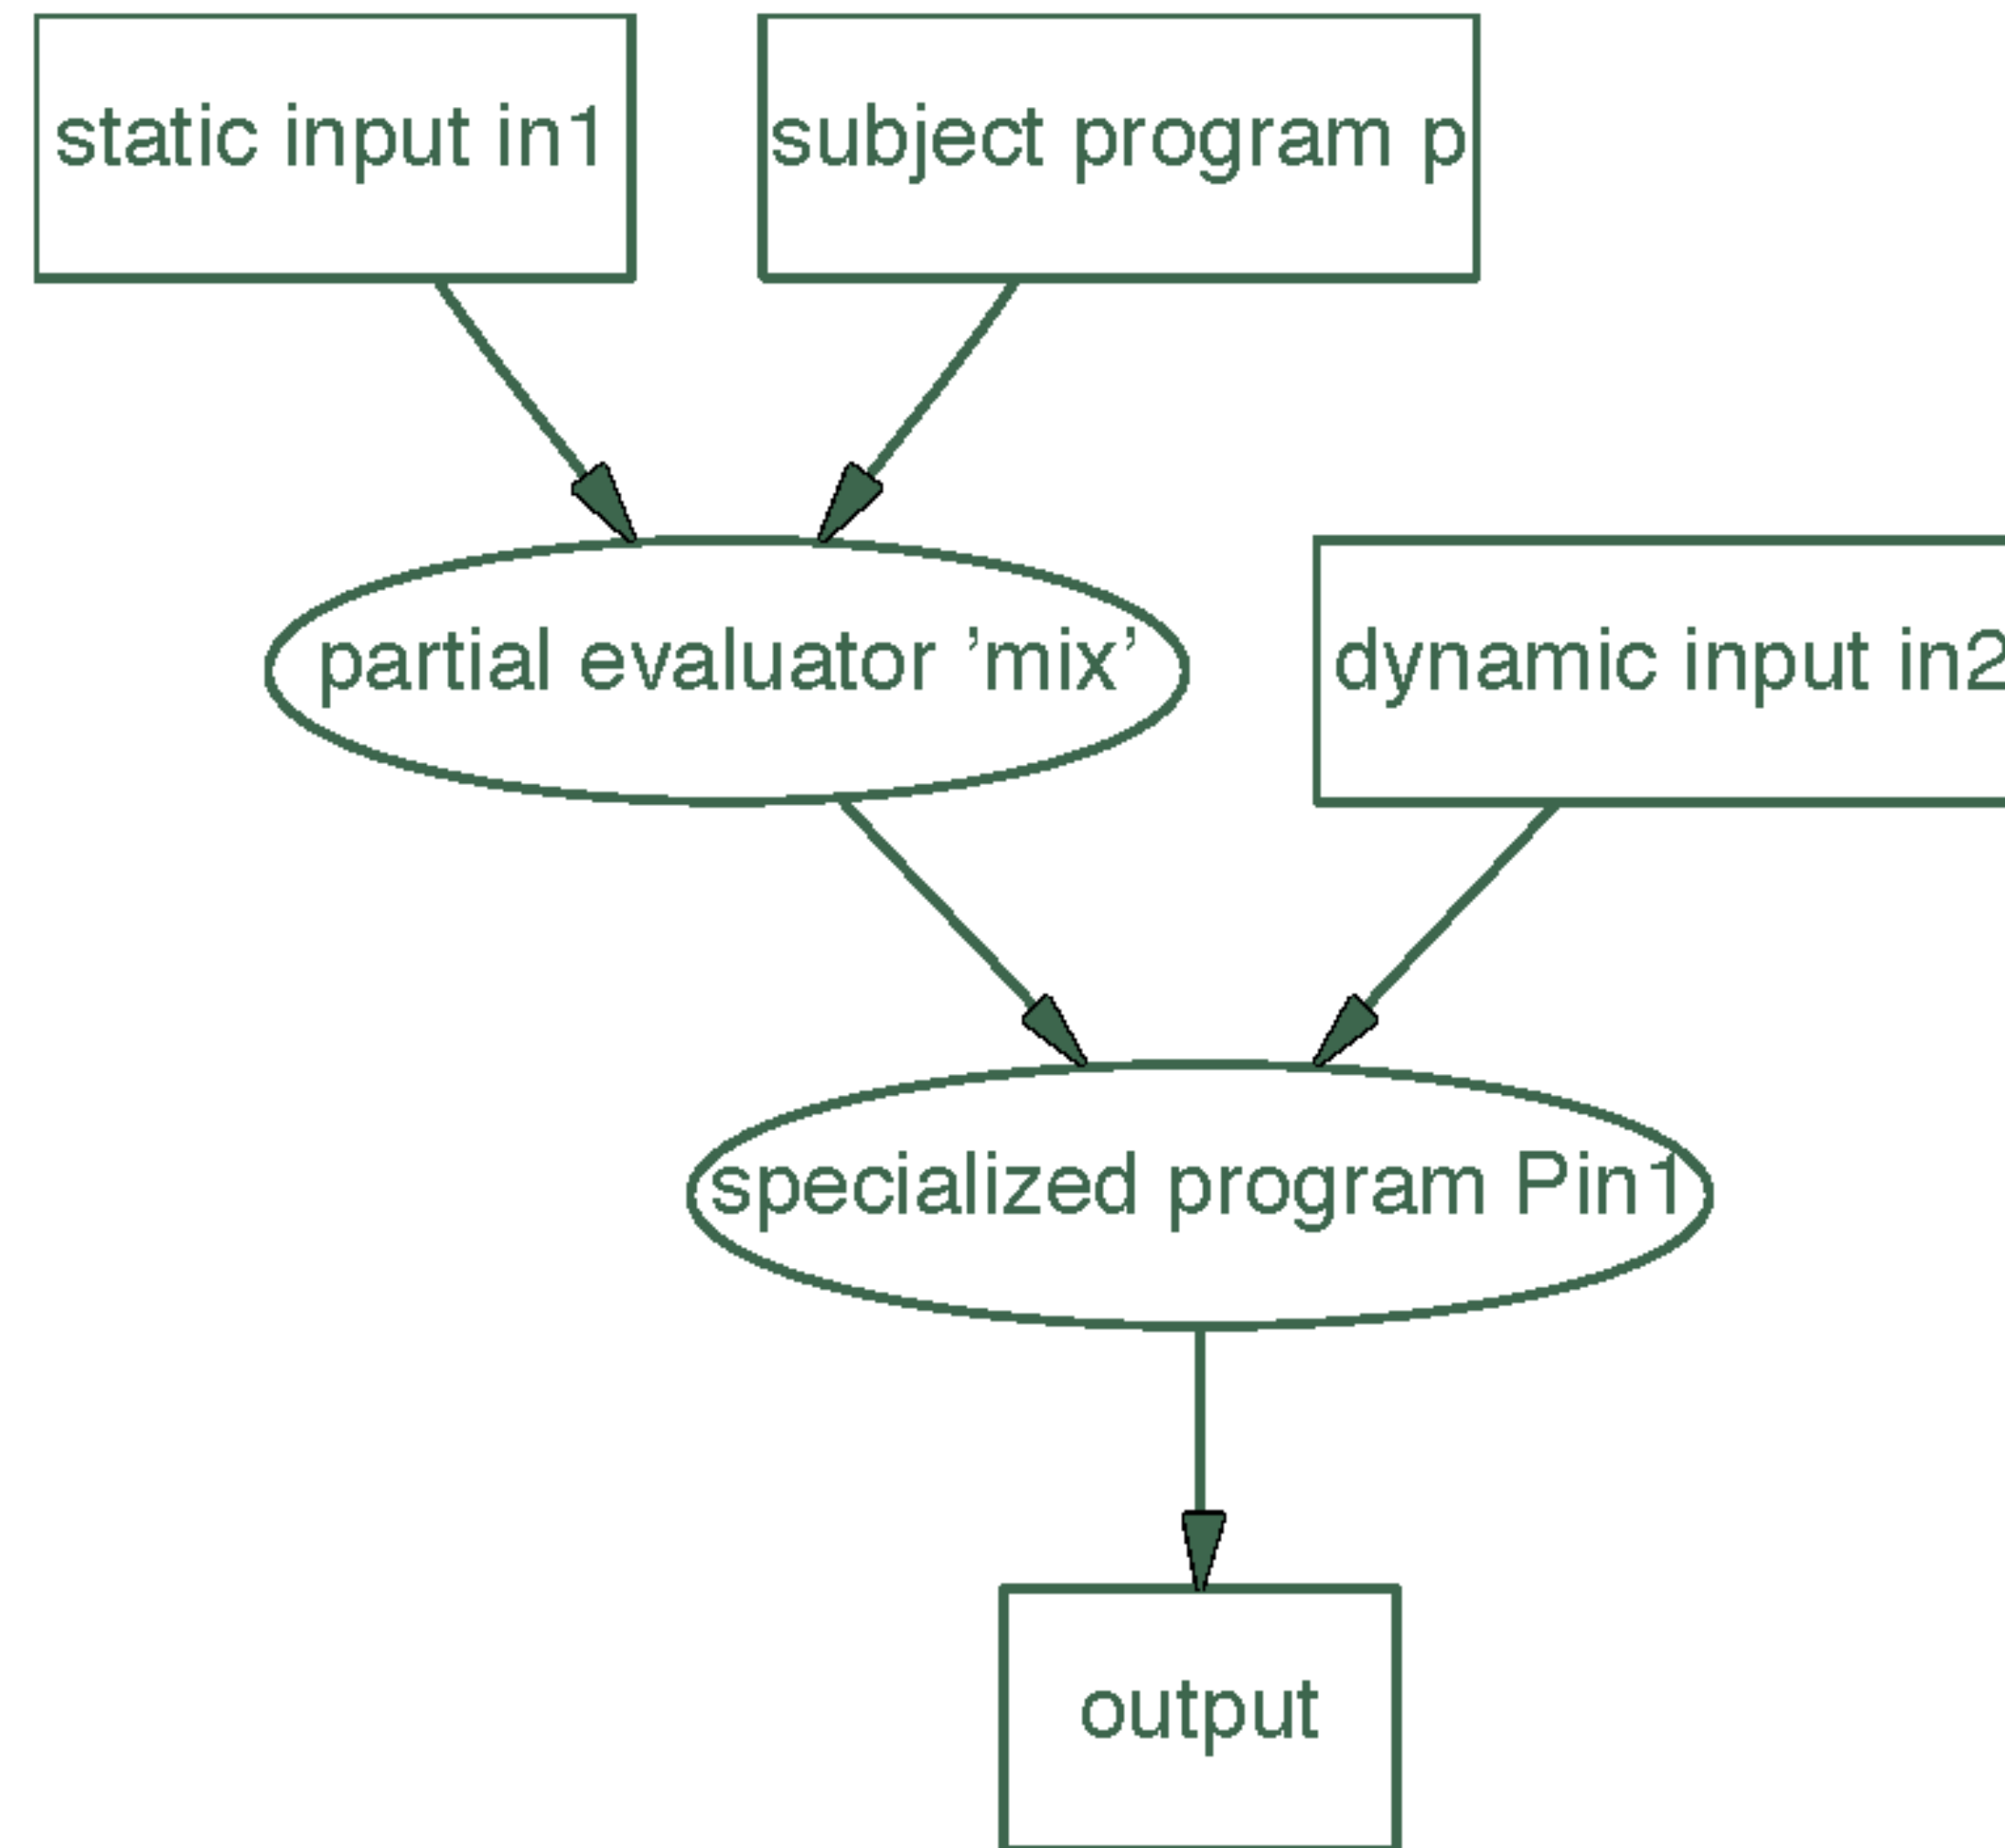
Source: Chapter 1 of Neil Jones, Carsten K. Gomar and Peter Sestoft. *Partial Evaluation and Automatic Program Generation*

Kleene's S-M-N Theorem

There is a primitive recursive function σ such that

$$f(x, y) = \sigma(f, x)(y)$$

Partial Evaluator (Mix)



Example

```
function power(x : int, n : int) : int =  
  if n = 0 then 1  
  else if even(n) then square(power(x, n/2))  
  else (x * power(x, n - 1))
```

Tiger

\Downarrow $n = 5$

Tiger

```
function power5(x : int) : int =  
  x * square(square(x))
```

precompute all expressions involving n
unfold recursive call to `power`
reduce $x*1$ to x

Techniques

- *Symbolic computation*
 - compute with expressions involving variables (not only values)
- *Unfolding*
 - replace call with instantiated body of function
- *Program point specialization*
 - create new version of function, specialized to some arguments

= definition + folding

Techniques

- *Definition*
 - introduce a new definition, or
 - extend an existing definition
- *Folding*
 - replace an expression with a function call
- *Memoization*
 - store the result of some computation
 - reproduce the result when computation is needed again

Terminology and Notation

- p : program
- $\llbracket p \rrbracket_L$: meaning of program p in language L
- L : implementation language
- S : source language
- T : target language
- $\llbracket - \rrbracket_L$ has type $D \rightarrow D^* \rightarrow D$

$$\text{output} = \llbracket p \rrbracket_L [\text{in}_1, \text{in}_2, \dots, \text{in}_n]$$

Equational Definition

one stage computation

$$\text{out} = \llbracket p \rrbracket [\text{in1}, \text{in2}]$$

two stage computation

$$p_{\text{in1}} = \llbracket \text{mix} \rrbracket [p, \text{in1}]$$

$$\text{out} = \llbracket p_{\text{in1}} \rrbracket \text{in2}$$

equational definition of `mix`

$$\llbracket p \rrbracket [\text{in1}, \text{in2}] = \llbracket \llbracket \text{mix} \rrbracket [p, \text{in1}] \rrbracket \text{in2}$$

different source, target, and implementation languages

$$\llbracket p \rrbracket_s [\text{in1}, \text{in2}] = \llbracket \llbracket \text{mix} \rrbracket_L [p, \text{in1}] \rrbracket_T \text{in2}$$

Why Partial Evaluation?

- Speedup

- $t_p(d_1, \dots, d_n)$: time to compute $\llbracket p \rrbracket_L [d_1, \dots, d_n]$
- $t_p(in1, in2)$: time to run program
- $t_{mix}(p, in1)$: time to specialize p to $in1$
- $t_{p_{in1}}(in2)$: time to run specialized program
- specialized program is faster (many applications to $in1$)

$$t_{p_{in1}}(in2) < t_p(in1, in2)$$

- specialization and running is faster (one application to $in1$)

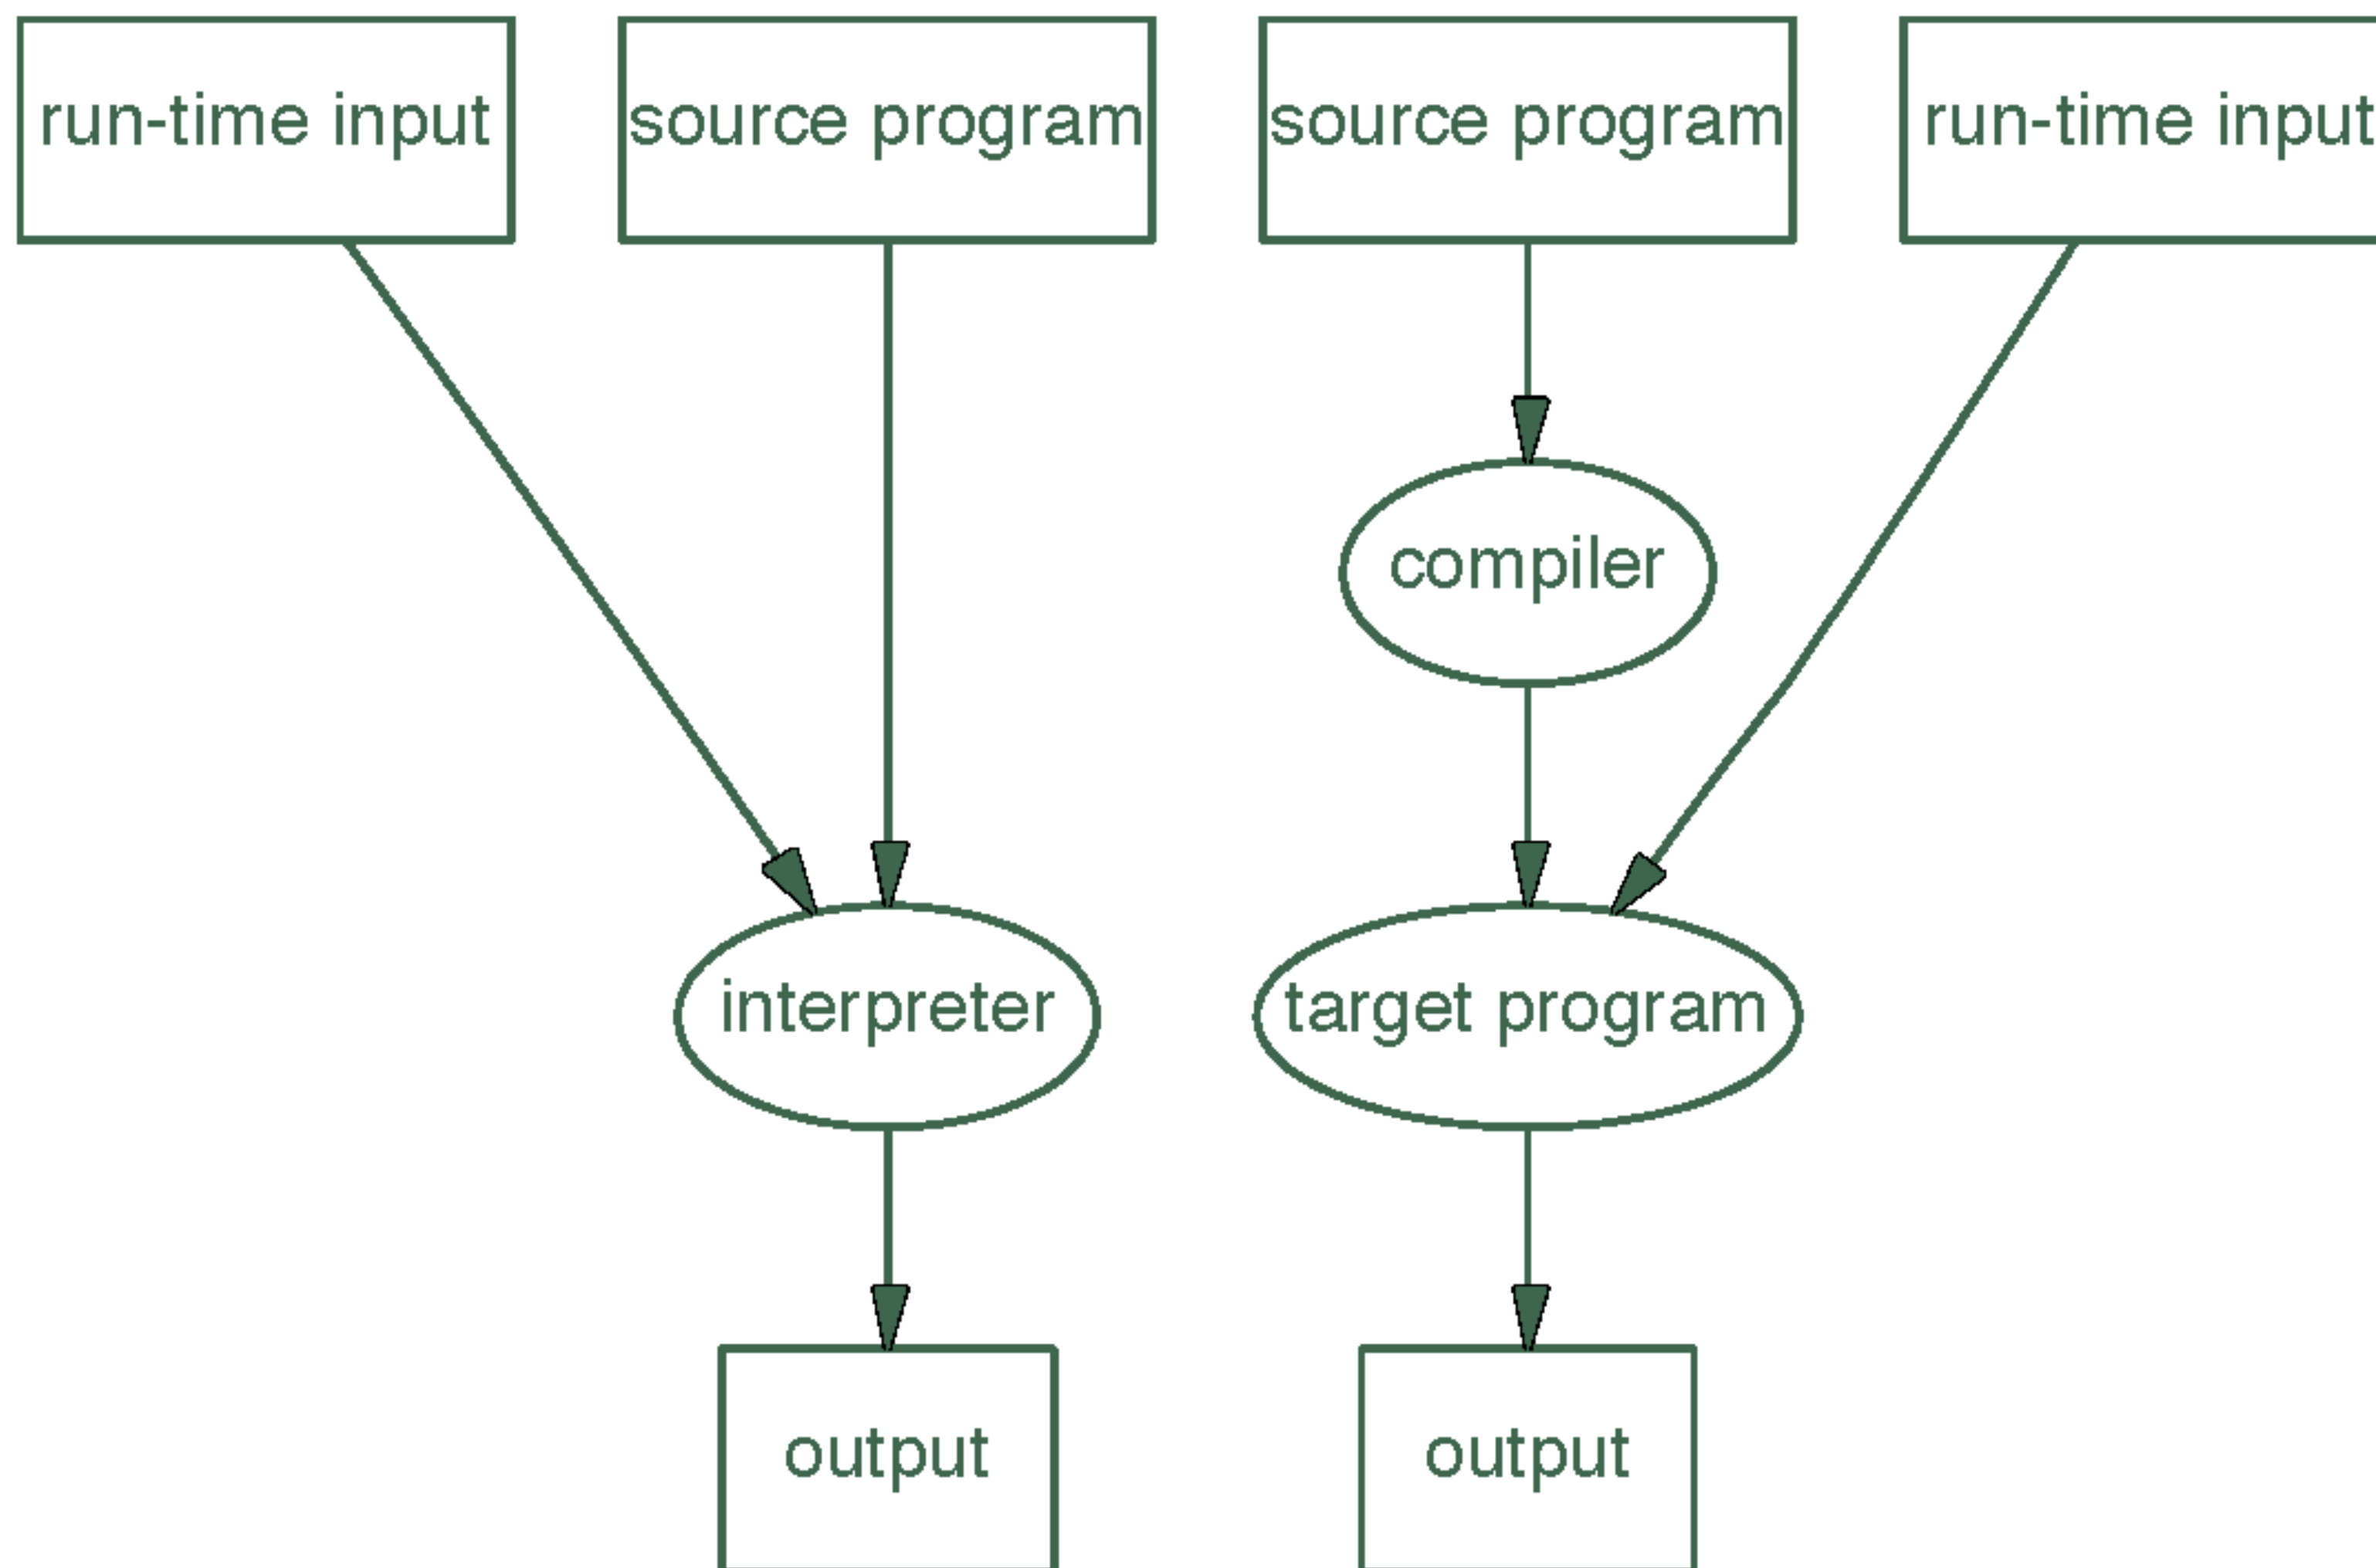
$$t_{mix}(p, in1) + t_{p_{in1}}(in2) < t_p(in1, in2)$$

Why Partial Evaluation?

- Efficient *and* modular solution
 - *program adapted to one problem*: efficient, but difficult to maintain
 - *modular/generic program*: easy to maintain and produce different instance, but inefficient
 - *partial evaluation*: specialize generic program to specific instance

How is this relevant for compilers and interpreters?

Staged Computation



Interpreters and Compilers

interpreter `int` executes a source program

$$\begin{aligned}\text{output} &= \llbracket \text{source} \rrbracket_S [\text{in}_1, \dots, \text{in}_n] \\ &= \llbracket \text{int} \rrbracket_L [\text{source}, \text{in}_1, \dots, \text{in}_n]\end{aligned}$$

compiler generates object program from source

$$\text{target} = \llbracket \text{compiler} \rrbracket_L [\text{source}]$$

running target

$$\begin{aligned}\text{output} &= \llbracket \text{source} \rrbracket_S [\text{in}_1, \dots, \text{in}_n] \\ &= \llbracket \text{target} \rrbracket_T [\text{in}_1, \dots, \text{in}_n]\end{aligned}$$

Parser Generators

generating a parser

$$\text{parser}_{\text{grammar}} = \llbracket \text{parsegen} \rrbracket_{\text{L}} [\text{grammar}]$$

parsing a string

$$\text{parsetree} = \llbracket \text{parser}_{\text{grammar}} \rrbracket_{\text{L}} [\text{string}]$$

general parsers

$$\text{parsetree} = \llbracket \text{genparser} \rrbracket_{\text{L}} [\text{grammar}, \text{string}]$$

Generalization: Executable Specifications

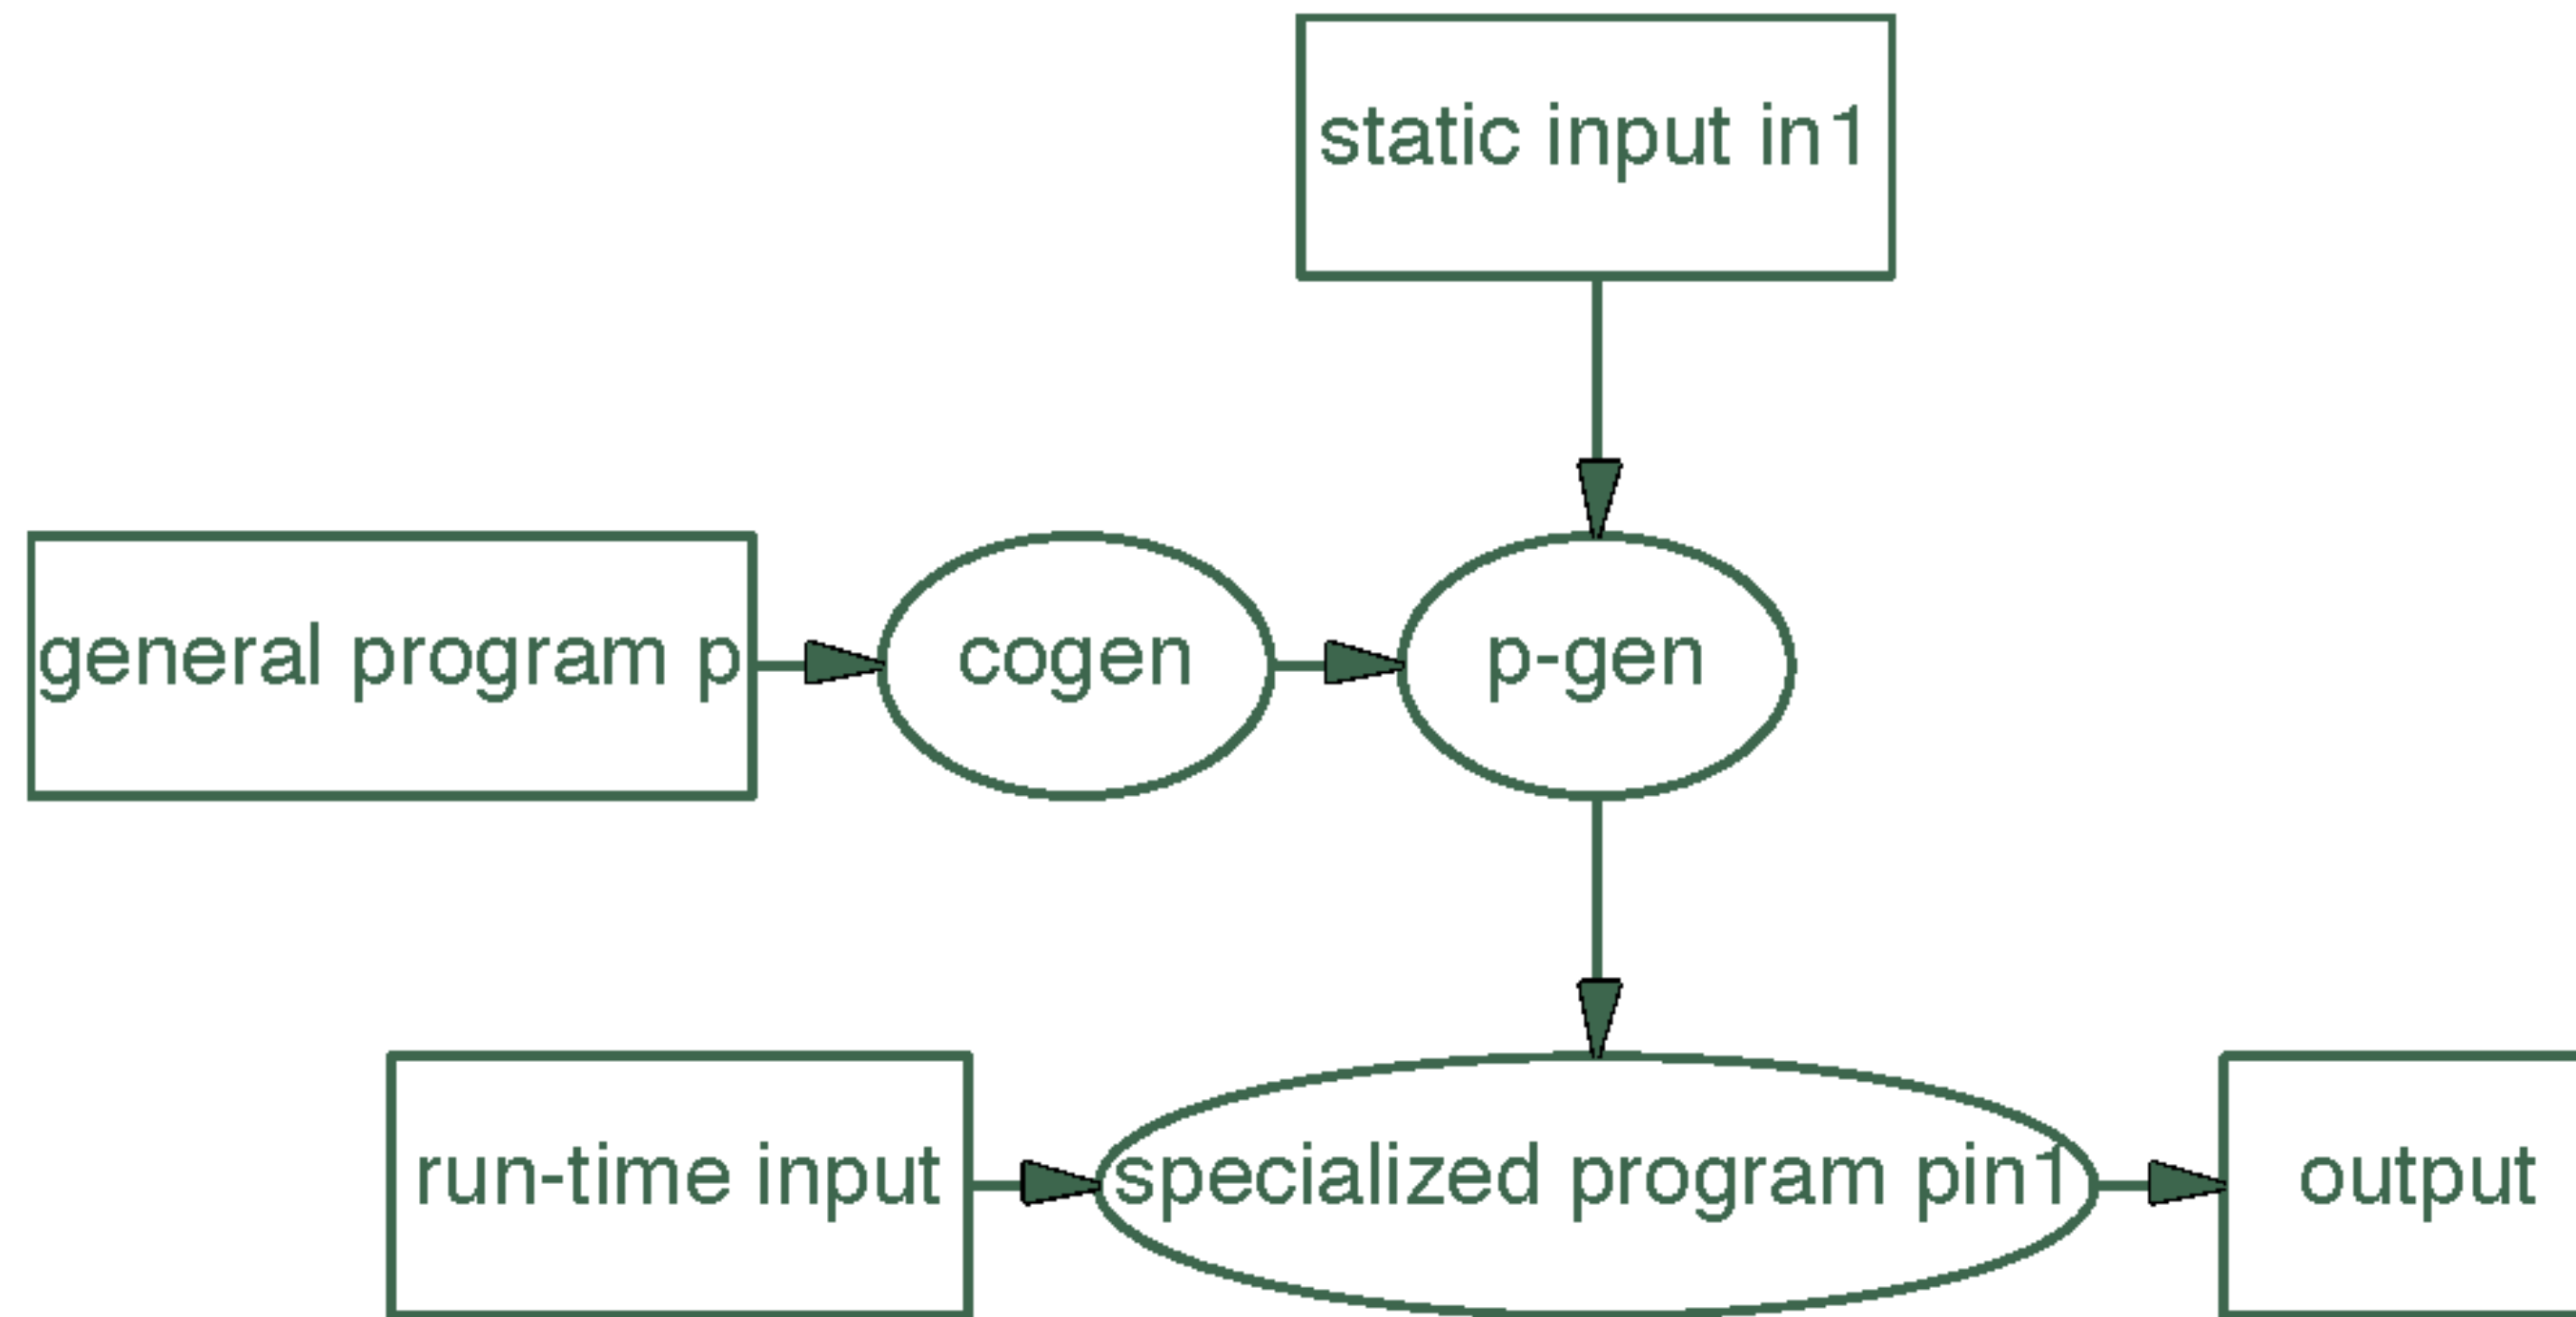
specification executer specexec 'runs' specification spec

$$\llbracket \text{specexec} \rrbracket [\text{spec}, \text{in}_1, \dots, \text{in}_n]$$

specexec is highly generic program

examples of such executers: int, genparser

Program Generator Generation



`compiler = [[cogen]] int`

`parsegen = [[cogen]] genparser`

Compiling by Partial Evaluation

First Futamura Projection (1971): specialize interpreter to source

$$\text{target} = \llbracket \text{mix} \rrbracket [\text{int}, \text{source}]$$

Compiling by partial evaluation yields correct target program:

$$\begin{aligned} \text{out} &= \llbracket \text{source} \rrbracket_s \text{input} \\ &= \llbracket \text{int} \rrbracket [\text{source}, \text{input}] \\ &= \llbracket \llbracket \text{mix} \rrbracket [\text{int}, \text{source}] \rrbracket \text{input} \\ &= \llbracket \text{target} \rrbracket \text{input} \end{aligned}$$

- target is in output language of mix
- specialization removes overhead of inspecting source

Compiler Generation

Second Futamura Projection

$\text{compiler} = \llbracket \text{mix} \rrbracket [\text{mix}, \text{int}]$

Correctness

$$\begin{aligned} \text{target} &= \llbracket \text{mix} \rrbracket [\text{int}, \text{source}] \\ &= \llbracket \llbracket \text{mix} \rrbracket [\text{mix}, \text{int}] \rrbracket \text{source} \\ &= \llbracket \text{compiler} \rrbracket \text{source} \end{aligned}$$

Compiler Generator Generation

Third Futamura Projection

$$\text{cogen} = \llbracket \text{mix} \rrbracket [\text{mix}, \text{mix}]$$

Correctness

$$\begin{aligned} \text{compiler} &= \llbracket \text{mix} \rrbracket [\text{mix}, \text{int}] \\ &= \llbracket \llbracket \text{mix} \rrbracket [\text{mix}, \text{mix}] \rrbracket \text{int} \\ &= \llbracket \text{cogen} \rrbracket \text{source} \end{aligned}$$

Partial Evaluation: Summary

Instance-specific program

- Can make very efficient
- Hard to maintain (performance interventions tangled with implementation)
- Hard to reuse for variants

Generic program

- Parametric in instances
- Easy to maintain
- Easy to reuse for variants
- Not efficient: overhead of parametricity

Partial evaluation

- Specialize generic program to specific instances

Some Research Projects

Specializing a meta-interpreter

- Meta-interpreter: DynSem specifications executed using interpreter
- Partial evaluation: remove two stages of interpretation
- Encouraging results using Truffle/Graal
- How close to native speed can we get?

Generate byte code compiler

- Input: DynSem specification
- Output: compiler from source to (custom) byte code
- Output: byte code interpreter

Partial evaluation

- Offline vs online partial evaluation
- Binding-time analysis

Multi-stage programming

- Explicit binding time annotations

Just-in-time (JIT) compilation

- Interpreter selectively compiles code (fragments) to machine code

Techniques for partial evaluation

- Meta-tracing
- Run-time specialisation (Truffle/Graal)

Except where otherwise noted, this work is licensed under

